

Automatic and Customizable Code Rewriting and Refactoring with Clang

Alistair Johnson

Johannes Doerfert

University of Oregon

Argonne National Lab

Disclaimer: This video reflects mid-October project status, changes may have since been made.

(Also presented at LLVM-HPC at SC21.)

Motivating Example: CUDA to HIP

```
cudaMalloc(&dev_array, num_bytes);
```

→

```
hipMalloc(&dev_array, num_bytes);
```

- Simple: search-and-replace `cudaMalloc` with `hipMalloc`

Motivating Example: CUDA to HIP

```
cudaMalloc(&dev_array, num_bytes);
```

→

```
hipMalloc(&dev_array, num_bytes);
```

- Simple: search-and-replace `cudaMalloc` with `hipMalloc`

```
custom_kernel<<<(N+255)/256, 256>>>(N, dev_array);
```

→

```
hipLaunchKernelGGL(custom_kernel, (N+255)/256, 256, 0, 0, N, dev_array);
```

- Complex: could possibly use `sed` or `awk`, but those are not optimal

Motivating Example: CUDA to HIP

```
[[clang::matcher("kernel_call")]]
auto cuda_kernel(int numBlocks, int numThreads) {
    auto arg1, arg2;
    {
        kernel<<<numBlocks, numThreads>>>(arg1, arg2);
    }
}
```

Motivating Example: CUDA to HIP

```
[[clang::matcher("kernel_call")]]
auto cuda_kernel(int numBlocks, int numThreads) {
    auto arg1, arg2;
    {
        kernel<<<numBlocks, numThreads>>>(arg1, arg2);
    }
}
```

```
[[clang::replace("kernel_call")]]
auto hip_kernel(int numBlocks, int numThreads) {
    auto arg1, arg2;
    {
        hipLaunchKernelGGL(kernel, numBlocks, numThreads, 0, 0, arg1, arg2);
    }
}
```

This Work

- A prototype (source-to-source) code rewriting tool built on Clang
- A user interface for this tool that is user-friendly and customizable
- An example implementation of `hipify` using this tool (in progress)

User Interface

- Pure C++ for simple learning curve
- New C++ attributes to define two sets of functions:
 - Matchers - describe which code to rewrite
 - `[[clang::matcher("<matcher_name>")]]`
 - Modifiers - describe how to rewrite matched code
 - `[[clang::replace("<list of matchers>")]]`
 - `[[clang::insert_before("<list of matchers>")]]`
 - `[[clang::insert_after("<list of matchers>")]]`
- Implemented with Clang's AST matcher and Rewriter interfaces

Use Cases

- API updates
- Switching APIs or programming models (e.g., CUDA to HIP)
- Adding (selective) instrumentation
- Device-specific code (e.g., OpenMP pragmas or kernels)
- Error-checking asserts
- Checkpointing
- Performance portability layers¹
- And many more!

¹J. K. Holmen, B. Peterson and M. Berzins, "An Approach for Indirectly Adopting a Performance Portability Layer in Large Legacy Codes," P3HPC 2019.