

Bringing up GlobalSel for optimized AArch64 codegen

Amara Emerson, Jessica Paquette - Apple

A quick recap...

What's GlobalSel?

- A new instruction selection framework, replacing FastISel & SelectionDAG
- Most notable difference is visibility of the entire function
 - SelectionDAG is limited to analyzing a single block at a time
- Other key design goals:
 - Customizability for targets to adapt to their specific needs
 - Compile time improvements

AArch64 GlobalSel backend

Project progress since COVID started

- GlobalSel was enabled for AArch64 at -O0 in 2018.
- Then attention shifted to the next big gigantic milestone: enabling at all optimization levels.
- Much more difficult than -O0, because:
 - IR input to codegen is extremely diverse
 - SelectionDAG is highly optimized
 - Small differences in performance/code size **matter**

Performance metrics


Performance metrics

How do we measure progress?

- Runtime performance (measured on Apple iOS devices)
 - SPEC2006, SPEC2017, Geekbench 5, test-suite, internal benchmarks
- Code size (-Os, -Oz)
 - CTMark, internal projects
- Compile time (-Os, -O3)
 - CTMark


Performance metrics

Runtime geomean regression vs SelectionDAG



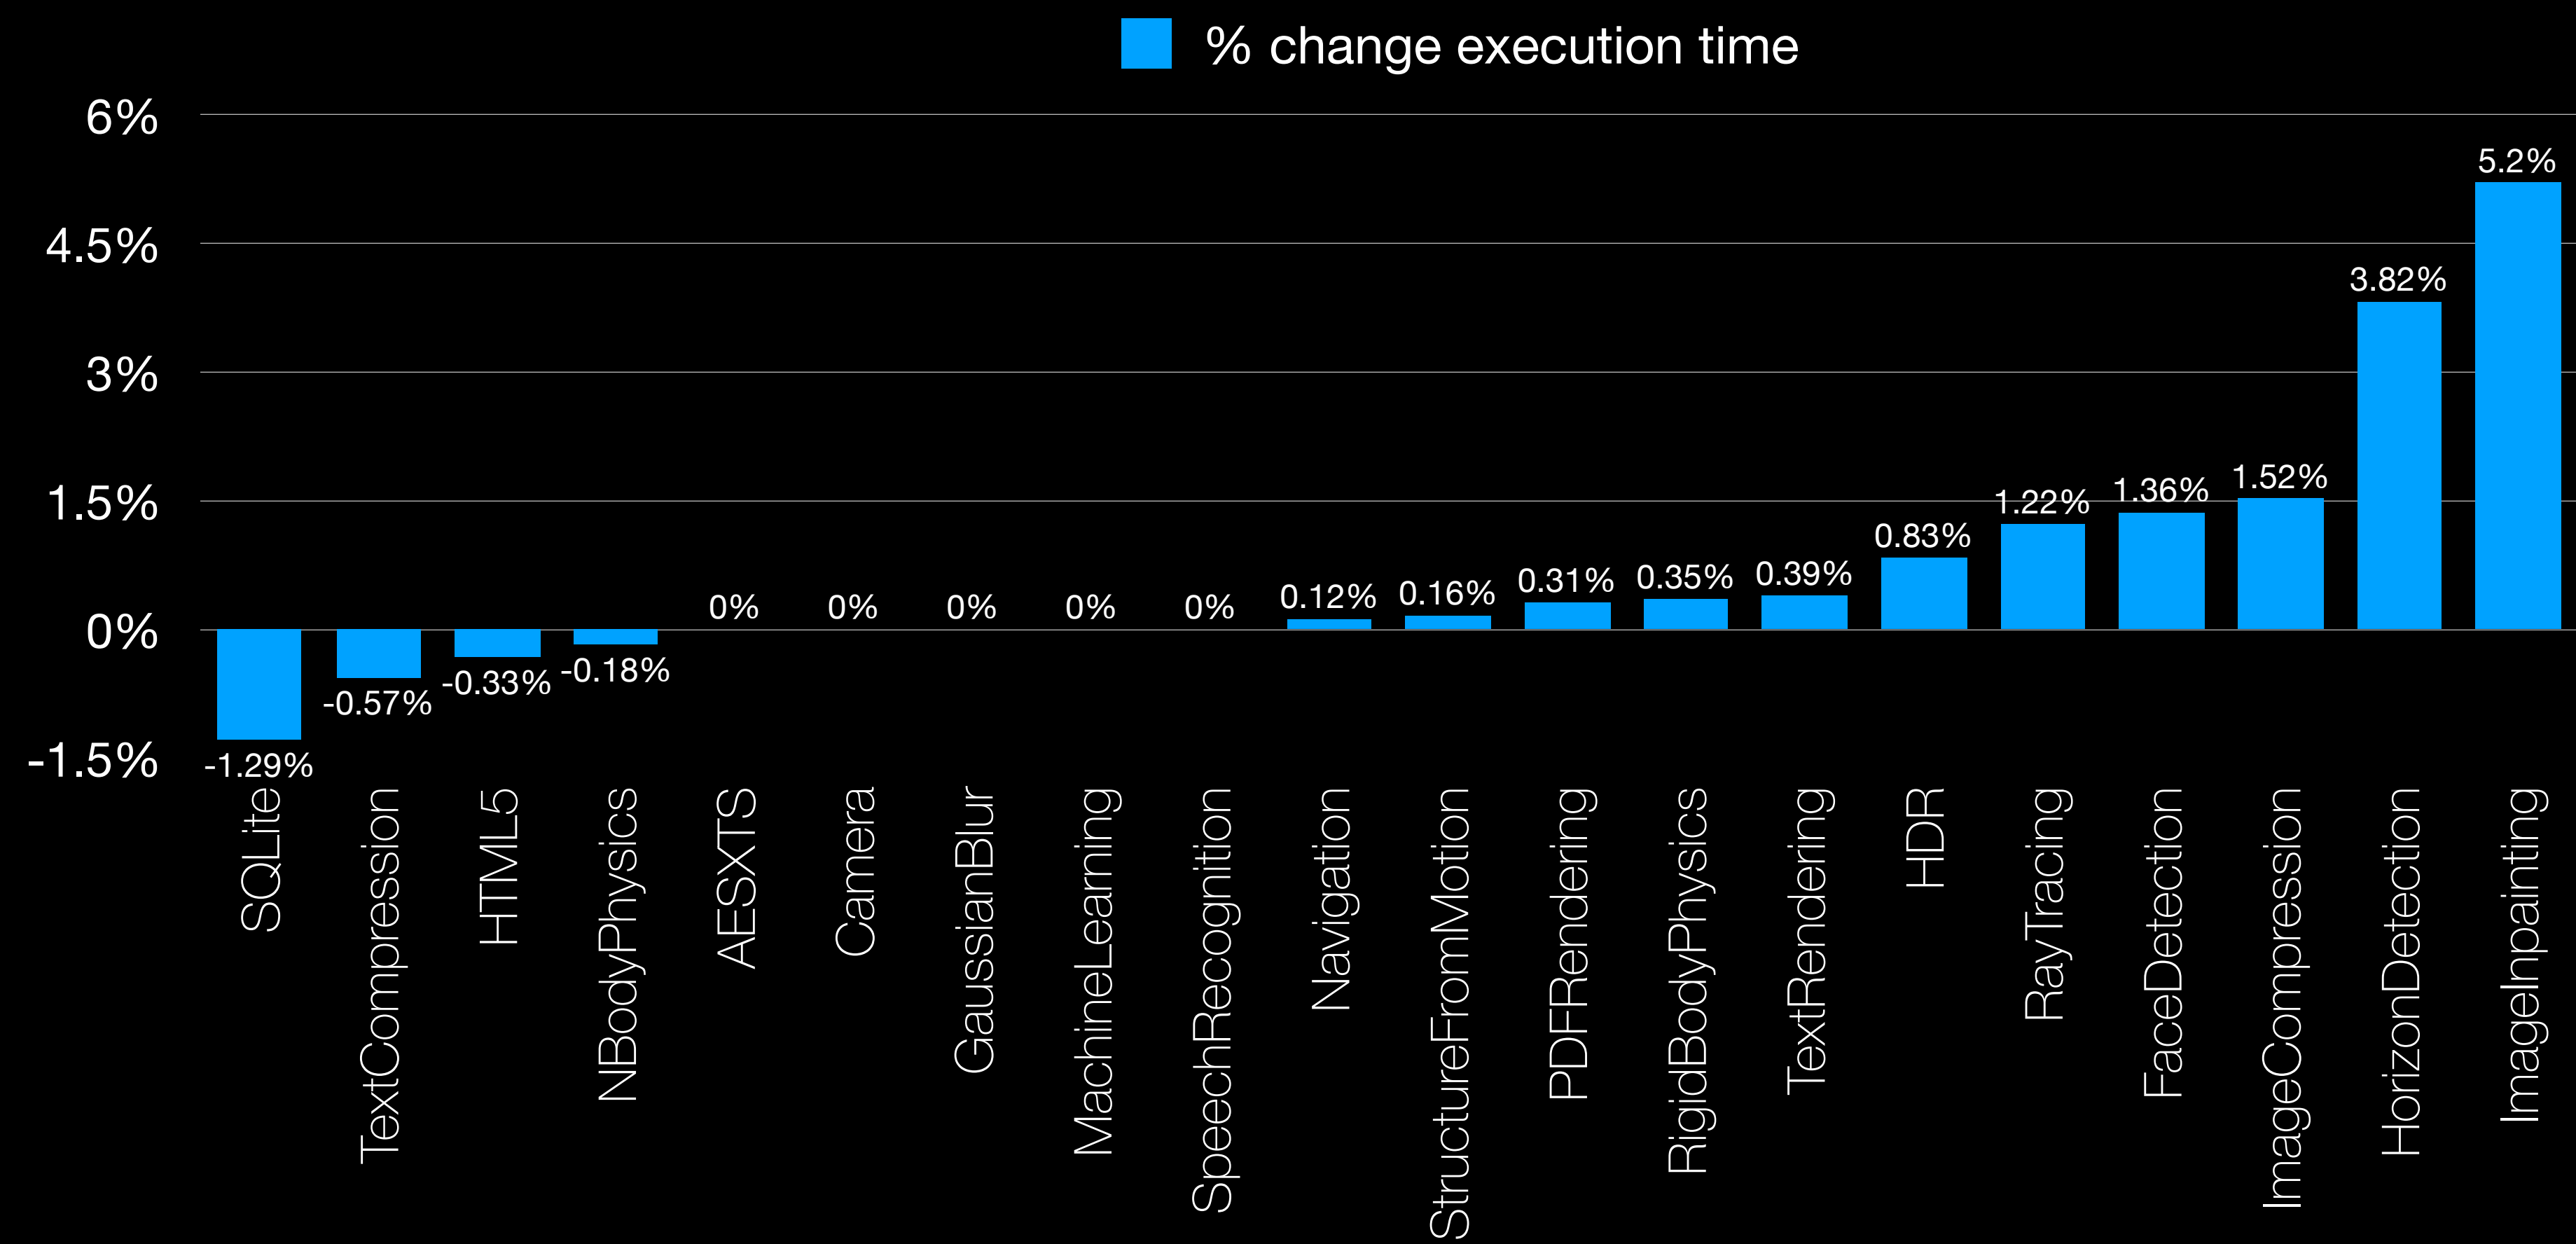
Performance metrics

Code size geomean regression vs SelectionDAG



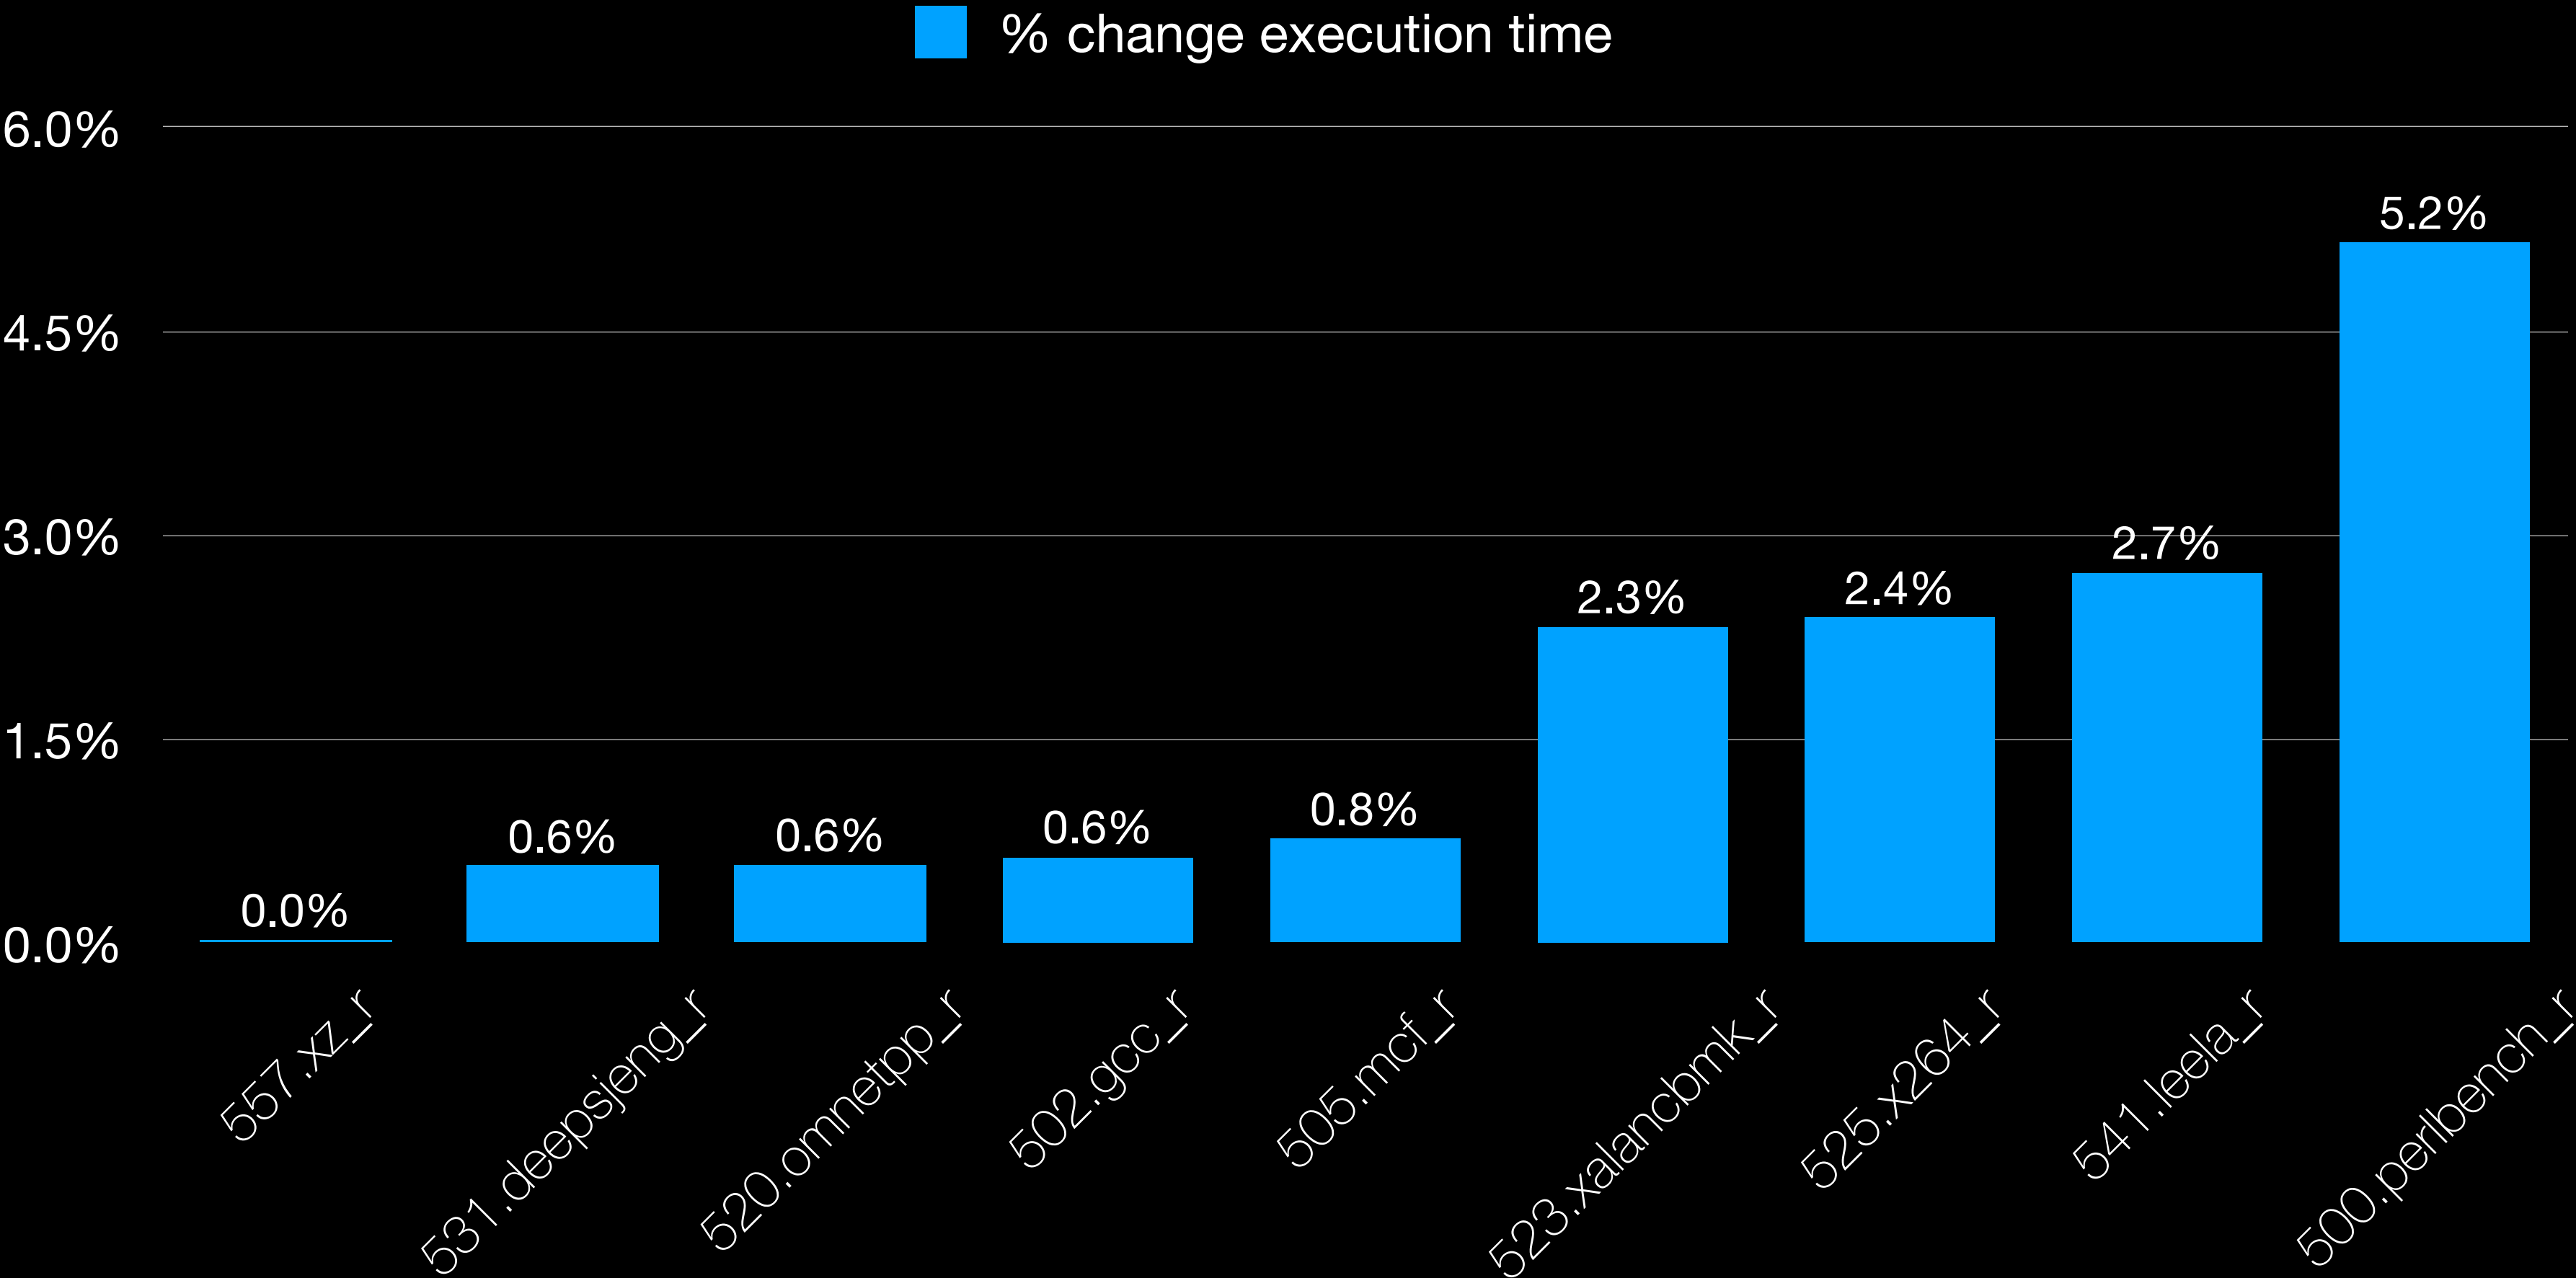
Performance metrics

Runtime geomean vs SelectionDAG: Geekbench - single threaded



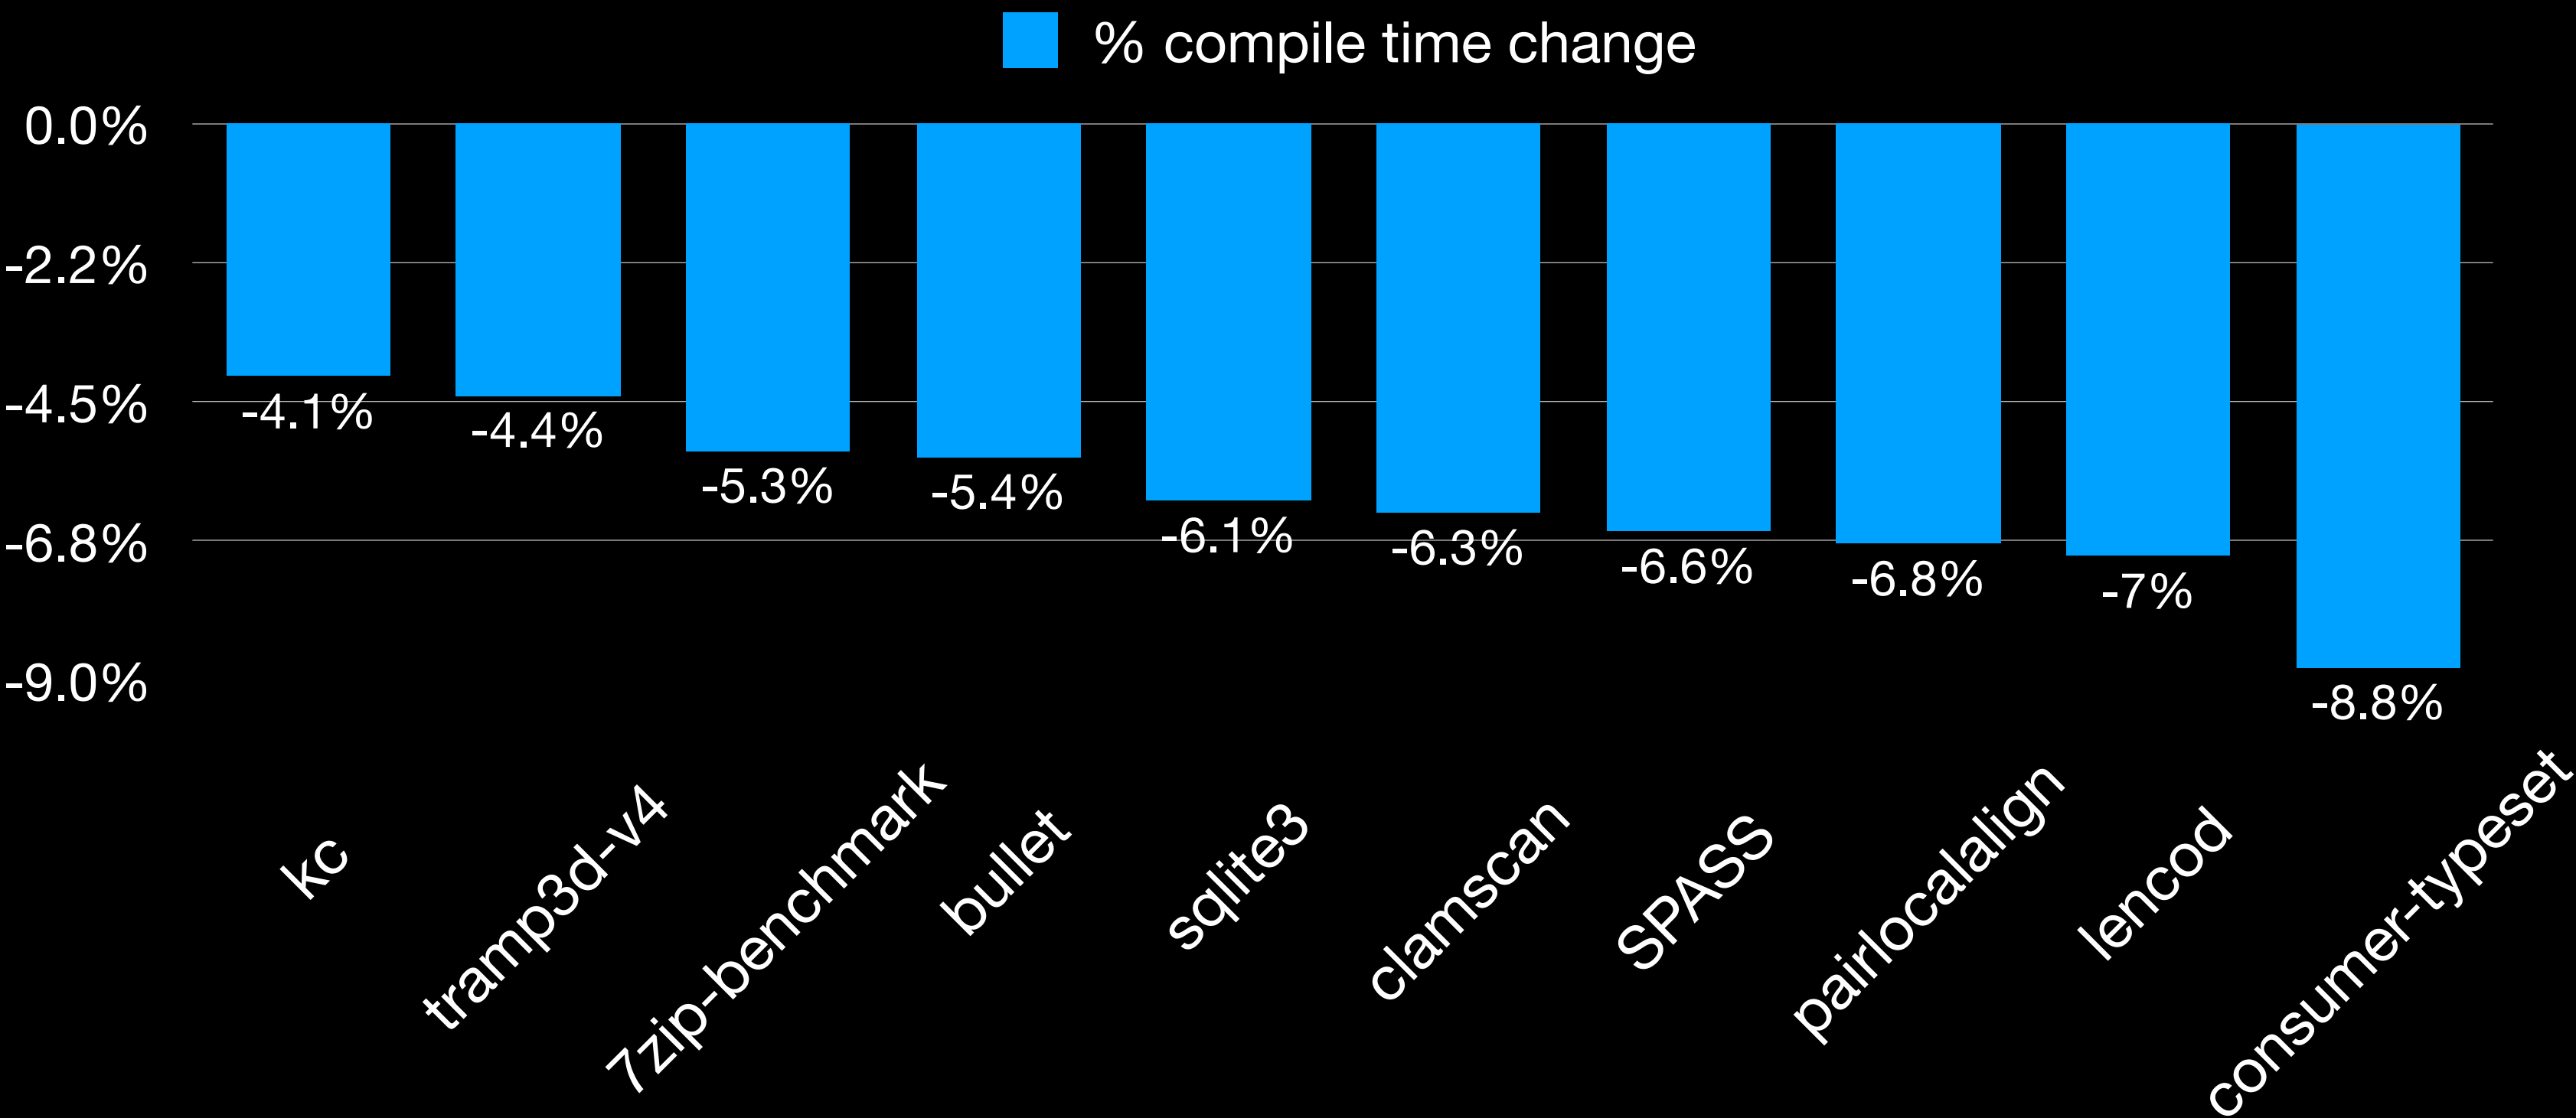
Performance metrics

Runtime geomean vs SelectionDAG: SPECINT2017 - INT



Performance metrics

ToT compile time vs SelectionDAG: CTMark -Os



Geomean improvement: **6.1%**


Performance metrics

SelectionDAG-only vs GlobalSel-only compile time: sqlite3

- Overall compile time improvement is nice, but front-end, IR opts, post-isel pipeline affect relative speedup.
- Isolating instruction selection only, how much faster is GlobalSel?
 - We include all combiners, legalizer code etc...

Performance metrics

SelectionDAG-only vs GlobalSel-only compile time: sqlite3 -Os



2.7x faster



Life of a GlobalSel instruction

`%x = add i64 %y, %z`



Translate IR to generic MIR

`%x:_(s64) = G_ADD %y, %z`



Legalize

`%x:_(s64) = G_ADD %y, %z`



Select register banks

`%x:gpr(s64) = G_ADD %y, %z`



Select target-specific MIR

`%x:gpr64(s64) = ADDSXrr %y, %z`

`%x = add i64 %y, %z`



Translate IR to generic MIR

`%x:_(s64) = G_ADD %y, %z`



Legalize

`%x:_(s64) = G_ADD %y, %z`



Select register banks

`%x:gpr(s64) = G_ADD %y, %z`



Select target-specific MIR

`%x:gpr64(s64) = ADDSXrr %y, %z`

IR Translation: translate LLVM IR
into gMIR instructions

`%x = add i64 %y, %z`

Translate IR to generic MIR

`%x:_(s64) = G_ADD %y, %z`

Legalize

`%x:_(s64) = G_ADD %y, %z`

Select register banks

`%x:gpr(s64) = G_ADD %y, %z`

Select target-specific MIR

`%x:gpr64(s64) = ADDSXrr %y, %z`

Legalization: mutate instruction using rules until it's "legal".

Legal in GlobalSel means that it's a valid, selectable instruction for the target.

`%x = add i64 %y, %z`



Translate IR to generic MIR

`%x:_(s64) = G_ADD %y, %z`



Legalize

`%x:_(s64) = G_ADD %y, %z`



Select register banks

`%x:gpr(s64) = G_ADD %y, %z`



Select target-specific MIR

`%x:gpr64(s64) = ADDSXrr %y, %z`

RegBankSelect: assigns register banks, as defined by the target, to minimize cost of cross-bank copies.

`%x = add i64 %y, %z`



Translate IR to generic MIR

`%x:_(s64) = G_ADD %y, %z`



Legalize

`%x:_(s64) = G_ADD %y, %z`



Select register banks

`%x:gpr(s64) = G_ADD %y, %z`



Select target-specific MIR

`%x:gpr64(s64) = ADDSXrr %y, %z`

Final instruction selection pass.
Matches patterns of gMIR into
target instructions.

AArch64 optimization pipeline

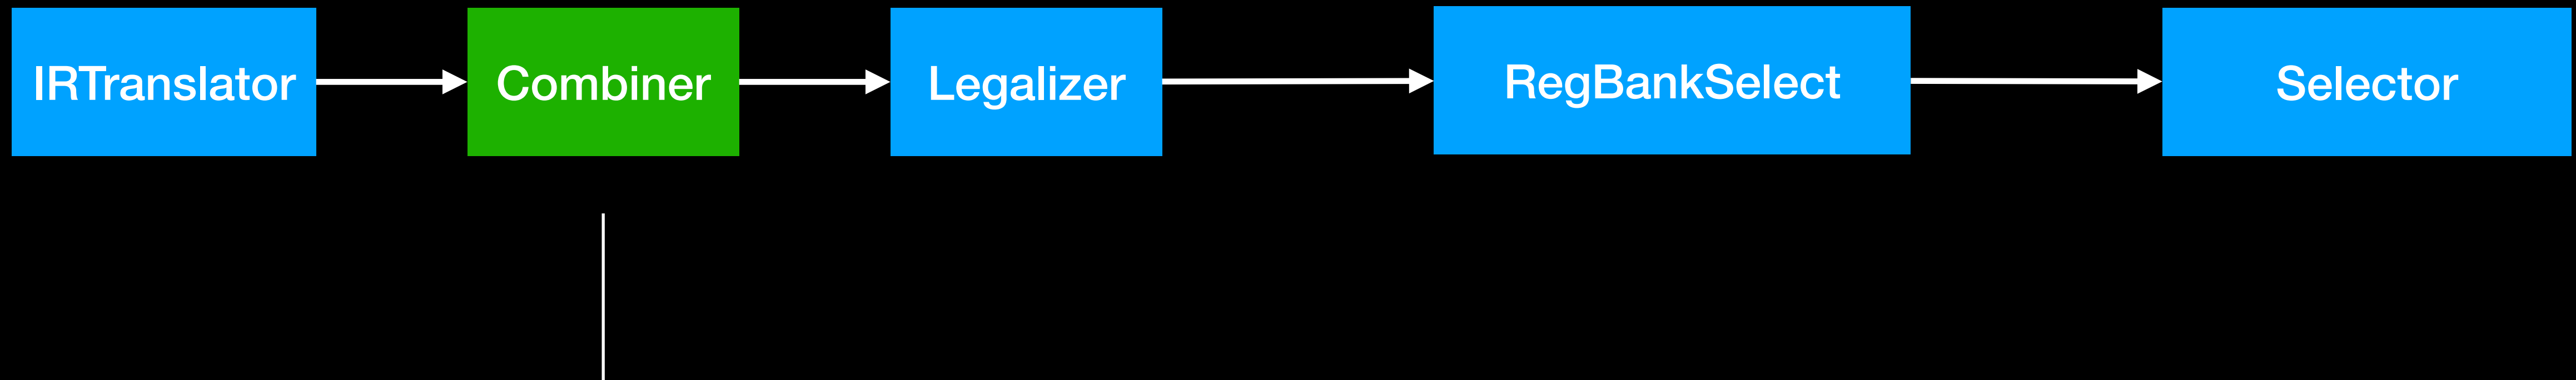
AArch64 optimization pipeline

Where should we place optimizations?



AArch64 optimization pipeline

The first combine pass



We call this the *Pre-Legalizer Combiner*

AArch64 optimization pipeline

The second combine pass



We call this the *Post-Legalizer Combiner*

AArch64 optimization pipeline

Selection-time optimization



Manual optimizations implemented at selection-time

Lowering complex operations

Lowering complex operations


A closer look at lowering and selection



- Some operations, like `G_SHUFFLE_VECTOR`, can be complex to optimize and select.
- For AArch64 we don't use custom legalization for shuffles.
- Let's take a closer look at the post-legalizer pipeline...

Lowering complex operations

A closer look at lowering and selection



- If a shuffle remains intact through the post-legalizer combiner, then the selector is responsible for all remaining optimization/lowering.
- This is possible, but adds a lot of complexity to one pass.
 - The selector cannot easily create new Generic MachineInstrs.
- We decided to add a lowering pass instead.

Lowering complex operations

A closer look at lowering and selection



- Post-legalizer lowering runs at all optimization levels.
 - It can optimize, but also contains necessary transforms for selection.
- We essentially split the selection process over two passes.
- Let's take a look at an example...

Lowering complex operations

Vector element reverse using G_SHUFFLE_VECTOR



- Our example is a G_SHUFFLE_VECTOR with a reverse element mask:

```
%shuf:_(<2 x s32>) = G_SHUFFLE_VECTOR %vec0(<2 x s32>), %vec1, shufflemask(1, 0)
```

```
%shuf:_(<2 x s32>) = G_SHUFFLE_VECTOR %vec0(<2 x s32>), %vec1, shufflemask(1, 0)
```

```
%shuf:_(<2 x s32>) = G_SHUFFLE_VECTOR %vec0(<2 x s32>), %vec1, shufflemask(1, 0)
```

```
def rev : GICombineRule<
  (defs root:$root, shuffle_matchdata:$matchinfo),
  (match (wip_match_opcode G_SHUFFLE_VECTOR):$root,
    [{ return matchREV(*${root}, MRI, ${matchinfo}); }]),
  (apply [{ applyShuffleVectorPseudo(*${root}, ${matchinfo}); }])
>;
```

```
%shuf:_(<2 x s32>) = G_SHUFFLE_VECTOR %vec0(<2 x s32>), %vec1, shufflemask(1, 0)
```

```
def rev : GICombineRule<
  (defs root:$root, shuffle_matchdata:$matchinfo),
  (match (wip_match_opcode G_SHUFFLE_VECTOR):$root,
    [{ return matchREV(*${root}, MRI, ${matchinfo}); }]),
  (apply [{ applyShuffleVectorPseudo(*${root}, ${matchinfo}); }])
>;
```



```
%shuf:_(<2 x s32>) = G_SHUFFLE_VECTOR %vec0(<2 x s32>), %vec1, shufflemask(1, 0)
```

```
static bool matchREV(MachineInstr &MI, MachineRegisterInfo &MRI,
                    ShuffleVectorPseudo &MatchInfo) {
    assert(MI.getOpcode() == TargetOpcode::G_SHUFFLE_VECTOR);
    ArrayRef<int> ShuffleMask = MI.getOperand(3).getShuffleMask();
    Register Dst = MI.getOperand(0).getReg();
    Register Src = MI.getOperand(1).getReg();
    LLT Ty = MRI.getType(Dst);
    unsigned EltSize = Ty.getScalarSizeInBits();
    // Element size for a rev cannot be 64.
    if (EltSize == 64)
        return false;
    unsigned NumElts = Ty.getNumElements();
    if (isREVMask(ShuffleMask, EltSize, NumElts, 64)) {
        MatchInfo = ShuffleVectorPseudo(AArch64::G_REV64, Dst, {Src});
        return true;
    }
    // ... G_REV32, G_REV64...
}
```

```
%shuf:_(<2 x s32>) = G_REV64 %vec0(<2 x s32>)
```

```
static bool matchREV(MachineInstr &MI, MachineRegisterInfo &MRI,
                    ShuffleVectorPseudo &MatchInfo) {
    assert(MI.getOpcode() == TargetOpcode::G_SHUFFLE_VECTOR);
    ArrayRef<int> ShuffleMask = MI.getOperand(3).getShuffleMask();
    Register Dst = MI.getOperand(0).getReg();
    Register Src = MI.getOperand(1).getReg();
    LLT Ty = MRI.getType(Dst);
    unsigned EltSize = Ty.getScalarSizeInBits();
    // Element size for a rev cannot be 64.
    if (EltSize == 64)
        return false;
    unsigned NumElt = Ty.getNumElements();
    if (isREVMask(ShuffleMask, EltSize, NumElt, 64)) {
        MatchInfo = ShuffleVectorPseudo(AArch64::G_REV64, Dst, {Src});
        return true;
    }
    // ... G_REV32, G_REV64...
}
```

`%shuf:_(<2 x s32>) = G_REV64 %vec0(<2 x s32>)`

```
def G_REV64 : AArch64GenericInstruction {  
  let OutOperandList = (outs type0:$dst);  
  let InOperandList = (ins type0:$src);  
  let hasSideEffects = 0;  
}
```

AArch64-only generic 64-bit rev

```
%shuf:_(<2 x s32>) = G_REV64 %vec0(<2 x s32>)
```

```
def G_REV64 : AArch64GenericInstruction {  
  let OutOperandList = (outs type0:$dst);  
  let InOperandList = (ins type0:$src);  
  let hasSideEffects = 0;  
}
```

AArch64-only generic 64-bit rev

```
def : GINodeEquiv<G_REV64, AArch64rev64>;
```

Import Tablegen patterns for AArch64rev64

```
%shuf:fpr(<2 x s32>) = G_REV64 %vec0(<2 x s32>)
```

```
%shuf:fpr(<2 x s32>) = G_REV64 %vec0(<2 x s32>)
```

```
def : GINodeEquiv<G_REV64, AArch64rev64>;
```

```
%shuf:fpr(<2 x s32>) = G_REV64 %vec0(<2 x s32>)
```

```
def : GINodeEquiv<G_REV64, AArch64rev64>;
```



Imported Tablegen rules for G_REV

```
(AArch64rev64:{*:[v2i32]} V64:{*:[v2i32]}:$Rn) => (REV64v2i32:{*:[v2i32]} V64:{*:[v2i32]}:$Rn)
(AArch64rev64:{*:[v2f32]} V64:{*:[v2f32]}:$Rn) => (REV64v2i32:{*:[v2f32]} V64:{*:[v2f32]}:$Rn)
(AArch64rev64:{*:[v4i16]} V64:{*:[v4i16]}:$Rn) => (REV64v4i16:{*:[v4i16]} V64:{*:[v4i16]}:$Rn)
...
```

```
%shuf:fpr(<2 x s32>) = G_REV64 %vec0(<2 x s32>)
```

```
def : GINodeEquiv<G_REV64, AArch64rev64>;
```



Imported Tablegen rules for G_REV

```
(AArch64rev64:{*:[v2i32]} V64:{*:[v2i32]}:$Rn) => (REV64v2i32:{*:[v2i32]} V64:{*:[v2i32]}:$Rn)
(AArch64rev64:{*:[v2f32]} V64:{*:[v2f32]}:$Rn) => (REV64v2i32:{*:[v2f32]} V64:{*:[v2f32]}:$Rn)
(AArch64rev64:{*:[v4i16]} V64:{*:[v4i16]}:$Rn) => (REV64v4i16:{*:[v4i16]} V64:{*:[v4i16]}:$Rn)
...
```



```
if (selectImpl(I, *CoverageInfo))
    return true;
```

selectImpl will try to select %shuf using an imported rule.


```
%shuf:fpr64 = REV64v2i32 %vec0
```

```
def : GINodeEquiv<G_REV64, AArch64rev64>;
```



Selected with 0 lines of code added to the selector!

Imported Tablegen rules for G_REV

```
(AArch64rev64:{*:[v2i32]} V64:{*:[v2i32]}:$Rn) => (REV64v2i32:{*:[v2i32]} V64:{*:[v2i32]}:$Rn)
(AArch64rev64:{*:[v2f32]} V64:{*:[v2f32]}:$Rn) => (REV64v2i32:{*:[v2f32]} V64:{*:[v2f32]}:$Rn)
(AArch64rev64:{*:[v4i16]} V64:{*:[v4i16]}:$Rn) => (REV64v4i16:{*:[v4i16]} V64:{*:[v4i16]}:$Rn)
...
```



```
if (selectImpl(I, *CoverageInfo))
    return true;
```

selectImpl will try to select %shuf using an imported rule.

The road ahead...

Remaining work

Addressing the long tail

- Fixing performance last major performance regressions is high priority.
- Improving general code size on -Os.
- Both of these have long tails of work. The last 1% is much harder than the first 10% of regressions to fix.
- ...but we're not far away. Expecting to be on par with SelectionDAG on the key metrics in the next few months.

Fallbacks

- When GlobalSel cannot handle some input IR, it “falls back” to SelectionDAG for that function.
- Across most test suites, average fallback rate is around 1% or less.
- Significant amounts of work went into reaching this level. Fallbacks now have a negligible effect on metrics.
- But eventual goal is to have no fallbacks at all.

Enabling?

- We think the benefits of GlobalSel in optimization opportunities and overall architecture are compelling.
- Plan to flip the switch for Darwin platforms when the metrics targets are reached.
- Eventually, we should show net improvements modulo some niche cases.

Acknowledgments

GlobalSel is a wide ranging effort

Special thanks to:

- Other targets, downstream and in tree, contributing heavily to development over the last few years.
- Occasional contributors for patches and reviews.
- Bug reports from external projects.

Q&A