

# Clang Static Analyzer: A Tryst with Smart Pointers

Deep Majumder

Indian Institute of Technology, Kharagpur



LLVM Developers' Meeting

November 17, 2021

# GSoC 2021

The talk revolves around my Google Summer of Code project this year under The LLVM Foundation. The project was **Making Smart Pointer Checkers default checkers in the Static Analyzer**.

# GSoC 2021

The talk revolves around my Google Summer of Code project this year under The LLVM Foundation. The project was **Making Smart Pointer Checkers default checkers in the Static Analyzer**.

My mentors were:

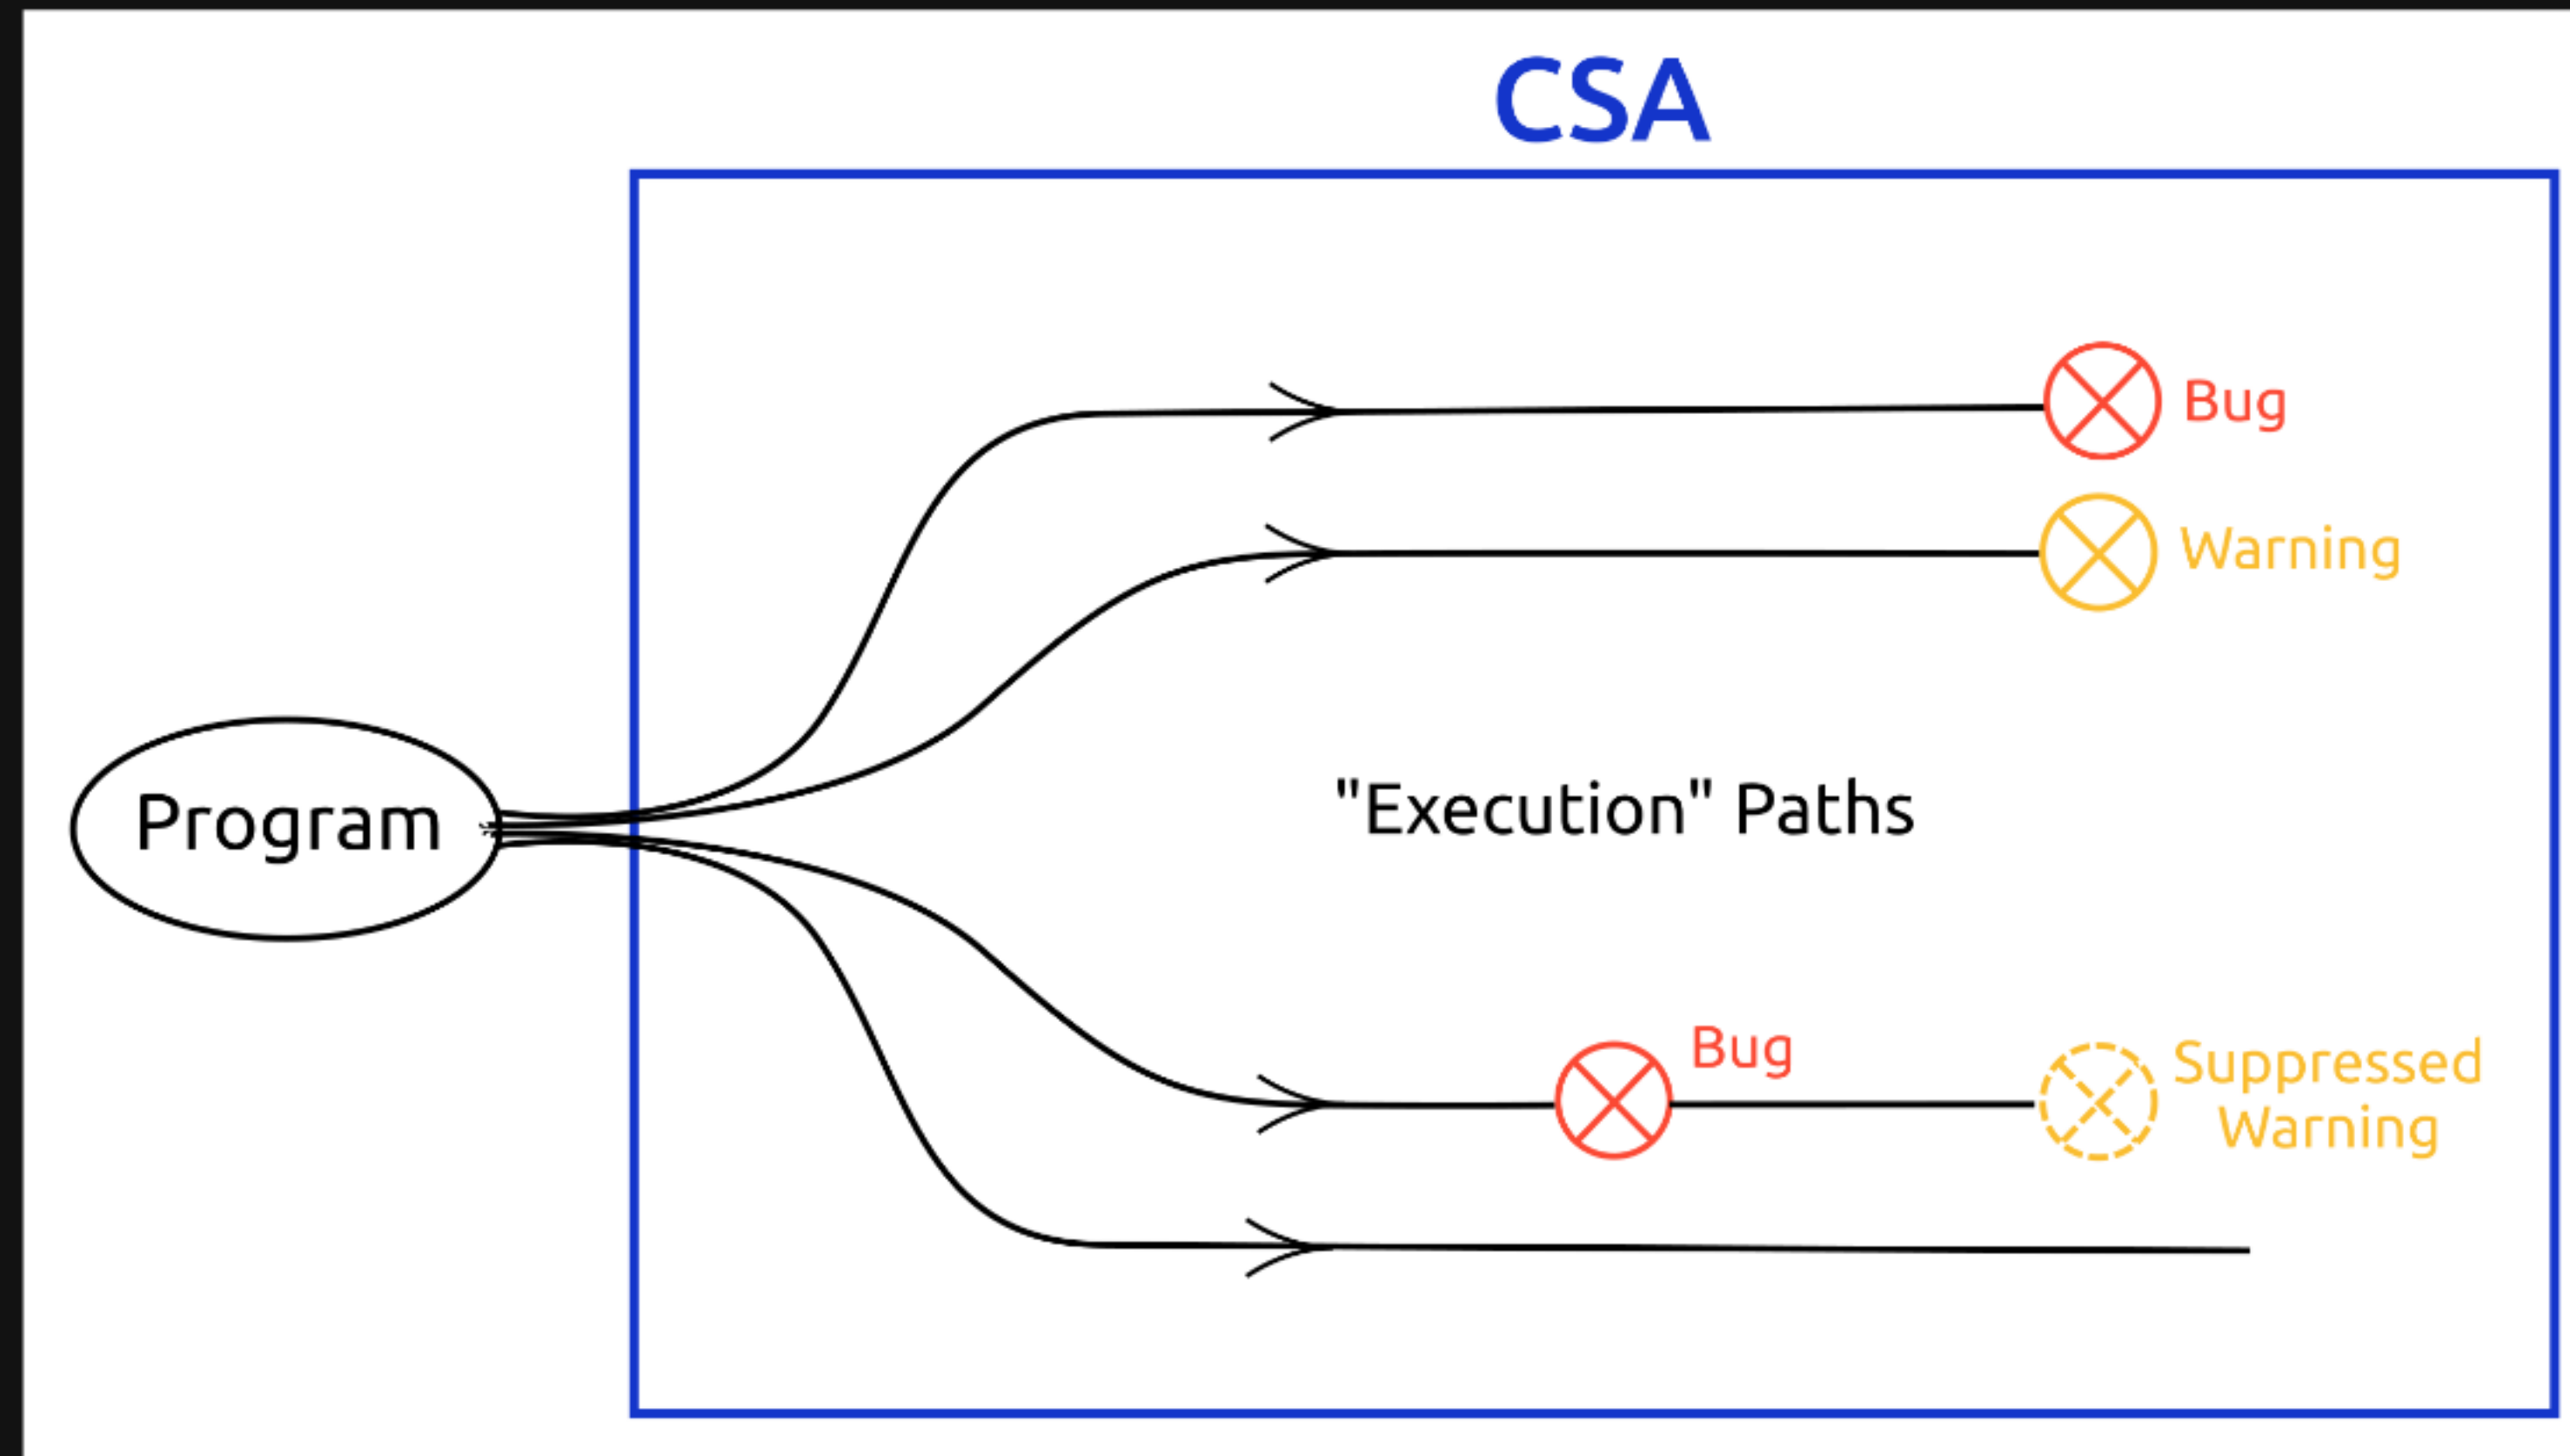
- Artem Dergachev (Apple Inc.)
- Gábor Horváth (Microsoft Corp.)
- Valeriy Savchenko (Apple Inc.)
- Raphael Iseemann (Vrije Universiteit Amsterdam)

# What is the Clang Static Analyzer?

The Clang Static Analyzer (CSA) is a static-analysis framework that uses **symbolic execution** over all possible paths in a program to flag various buggy scenarios.

# What is the Clang Static Analyzer?

The Clang Static Analyzer (CSA) is a static-analysis framework that uses **symbolic execution** over all possible paths in a program to flag various buggy scenarios.



The CSA tries to have no (low) false positives. This means:

- A bug reported by the CSA is **definitely** a bug.
- Absence of bug reports by the CSA does **not** indicate absence of bugs.

This is different from the model followed by some other static analysis systems, such as the one in Rust.



# It's Checkers all the way down!

Checkers are the backbone of the CSA.

# It's Checkers all the way down!

Checkers are the backbone of the CSA.

Checkers can:

- Create data
- Modify existing data
- Use the data to detect bugs



# It's Checkers all the way down!

Checkers are the backbone of the CSA.

Checkers can:

- Create data
- Modify existing data
- Use the data to detect bugs

Roughly speaking, **data** stands for constraints on the possible values of a variable (or expression) at some point in the program. Also, we can store **metadata** to help dealing with complex scenarios.

For dealing with smart-pointers, we have currently have two checkers:

1. `SmartPointerModelling`, which actually creates the requisite data
2. `SmartPointerChecker`, which uses the data to detect null-dereferences

This de-coupling allows us to have **multiple** checkers to model the various smart-pointers, while having a *single checker emit the bug reports*.

# How is it modelled?

`SmartPtrModelling` models the behaviour of `std::unique_ptr` by keeping track of the *nullity of the raw-pointer inside the smart-pointer*.

```
void foo() {
    unique_ptr<int> ptr = make_unique<int>(13);
    // ptr definitely is not null now.
    cout << *ptr << "\n";
    // This statement is safe.
    ptr.reset();
    // ptr now is definitely null.
    magicFunc(*ptr);
    // This statement is definitely a
    // null-ptr dereference.
}
```

```
void bar(unique_ptr<int> ptr) {
    // At this point we know nothing about ptr
    *ptr;
    // Since we aren't sure, we emit no report here
    if (!ptr) {
        // We are now sure that ptr is null here
        magicFunc(*ptr);
        // This statement is definitely a
        // null-ptr dereference.
    }
}
```

# How is it modelled?

`SmartPtrModelling` models the behaviour of `std::unique_ptr` by keeping track of the *nullity of the raw-pointer inside the smart-pointer*.

```
void foo() {  
    unique_ptr<int> ptr = make_unique<int>(13);  
    // ptr definitely is not null now.  
    cout << *ptr << "\n";  
    // This statement is safe.  
    ptr.reset();  
    // ptr now is definitely null.  
    magicFunc(*ptr);  
    // This statement is definitely a  
    // null-ptr dereference.  
}
```

```
void bar(unique_ptr<int> ptr) {  
    // At this point we know nothing about ptr  
    *ptr;  
    // Since we aren't sure, we emit no report here  
    if (!ptr) {  
        // We are now sure that ptr is null here  
        magicFunc(*ptr);  
        // This statement is definitely a  
        // null-ptr dereference.  
    }  
}
```



# How is it modelled?

`SmartPtrModelling` models the behaviour of `std::unique_ptr` by keeping track of the *nullity of the raw-pointer inside the smart-pointer*.

```
void foo() {
    unique_ptr<int> ptr = make_unique<int>(13);
    // ptr definitely is not null now.
    cout << *ptr << "\n";
    // This statement is safe.
    ptr.reset();
    // ptr now is definitely null.
    magicFunc(*ptr);
    // This statement is definitely a
    // null-ptr dereference.
}
```

```
void bar(unique_ptr<int> ptr) {
    // At this point we know nothing about ptr
    *ptr;
    // Since we aren't sure, we emit no report here
    if (!ptr) {
        // We are now sure that ptr is null here
        magicFunc(*ptr);
        // This statement is definitely a
        // null-ptr dereference.
    }
}
```

# How is it modelled?

`SmartPtrModelling` models the behaviour of `std::unique_ptr` by keeping track of the *nullity of the raw-pointer inside the smart-pointer*.

```
void foo() {
    unique_ptr<int> ptr = make_unique<int>(13);
    // ptr definitely is not null now.
    cout << *ptr << "\n";
    // This statement is safe.
    ptr.reset();
    // ptr now is definitely null.
    magicFunc(*ptr);
    // This statement is definitely a
    // null-ptr dereference.
}
```

```
void bar(unique_ptr<int> ptr) {
    // At this point we know nothing about ptr
    *ptr;
    // Since we aren't sure, we emit no report here
    if (!ptr) {
        // We are now sure that ptr is null here
        magicFunc(*ptr);
        // This statement is definitely a
        // null-ptr dereference.
    }
}
```

# How is it modelled?

`SmartPtrModelling` models the behaviour of `std::unique_ptr` by keeping track of the *nullity of the raw-pointer inside the smart-pointer*.

```
void foo() {
    unique_ptr<int> ptr = make_unique<int>(13);
    // ptr definitely is not null now.
    cout << *ptr << "\n";
    // This statement is safe.
    ptr.reset();
    // ptr now is definitely null.
    magicFunc(*ptr);
    // This statement is definitely a
    // null-ptr dereference.
}
```

```
void bar(unique_ptr<int> ptr) {
    // At this point we know nothing about ptr
    *ptr;
    // Since we aren't sure, we emit no report here
    if (!ptr) {
        // We are now sure that ptr is null here
        magicFunc(*ptr);
        // This statement is definitely a
        // null-ptr dereference.
    }
}
```



# How is it modelled?

`SmartPtrModelling` models the behaviour of `std::unique_ptr` by keeping track of the *nullity of the raw-pointer inside the smart-pointer*.

```
void foo() {  
    unique_ptr<int> ptr = make_unique<int>(13);  
    // ptr definitely is not null now.  
    cout << *ptr << "\n";  
    // This statement is safe.  
    ptr.reset();  
    // ptr now is definitely null.  
    magicFunc(*ptr);  
    // This statement is definitely a  
    // null-ptr dereference.  
}
```

```
void bar(unique_ptr<int> ptr) {  
    // At this point we know nothing about ptr  
    *ptr;  
    // Since we aren't sure, we emit no report here  
    if (!ptr) {  
        // We are now sure that ptr is null here  
        magicFunc(*ptr);  
        // This statement is definitely a  
        // null-ptr dereference.  
    }  
}
```

What is there to model?

Everything!

# What is there to model?

## Everything!

When we are modelling a C++ class like `std::unique_ptr`, we need to model all aspects of it. This includes modelling at least all **member** functions, constructors, and destructors.

# What is there to model?

## Everything!

When we are modelling a C++ class like `std::unique_ptr`, we need to model all aspects of it. This includes modelling at least all **member** functions, **constructors**, and **destructors**.

The aim of modelling a class in a checker is to **add** information which is specific to the semantics of the class at hand.

# What is there to model?

## Everything!

When we are modelling a C++ class like `std::unique_ptr`, we need to model all aspects of it. This includes modelling at least all **member** functions, **constructors**, and **destructors**.

The aim of modelling a class in a checker is to **add** information which is specific to the semantics of the class at hand.

This information is not necessarily discoverable automatically because the source code **may not be available**. After all, the interface to a C++ library is via the header file, which may not contain all the code.



# The Perils of Incomplete Modelling

**Loss of information:** Suppose we don't model the `reset()` method for `std::unique_ptr`, and assume its source code is not available.

CSA will default to **conservative evaluation**, which basically **removes** information which may have been modified by the modelled method.

# The Perils of Incomplete Modelling

**Loss of information:** Suppose we don't model the `reset()` method for `std::unique_ptr`, and assume its source code is not available.

CSA will default to **conservative evaluation**, which basically **removes** information which may have been modified by the modelled method.

Then, the following code will report no errors, when there is obviously one (a *false-negative*).

```
void foo() {
    unique_ptr<string> strPtr = make_unique<string>("Oh no!");
    strPtr.reset();
    auto len = strPtr->size(); // This is null-pointer dereference!
    // But we have no warning!
}
```



**Inconsistent information:** Suppose we don't model the `destructor` of `std::unique_ptr`.

Since the code for the destructor is usually available, CSA will automatically evaluate it, in a process known as **inlining**.

**Inconsistent information:** Suppose we don't model the `~destructor` of `std::unique_ptr`.

Since the code for the destructor is usually available, CSA will automatically evaluate it, in a process known as **inlining**.

The destructor for `std::unique_ptr` is (very) roughly as follows:

```
~unique_ptr() {  
    auto &rawPtr = innerPtr.getPtr();  
    if (rawPtr != nullptr)  
        get_deleter()(std::move(rawPtr));  
    // This defaults to `delete rawPtr;`  
}
```

**Inconsistent information:** Suppose we don't model the `~destructor` of `std::unique_ptr`.

Since the code for the destructor is usually available, CSA will automatically evaluate it, in a process known as **inlining**.

The destructor for `std::unique_ptr` is (very) roughly as follows:

```
~unique_ptr() {  
    auto &rawPtr = innerPtr.getPtr();  
    if (rawPtr != nullptr)  
        get_deleter()(std::move(rawPtr));  
    // This defaults to `delete rawPtr;`  
}
```

Sometimes "bug" reports generated by the CSA are **suppressed** when we believe that they are probably false-positives. "Bugs" that come from the standard library are one such class.

But as a side-effect, the following code's leak warning (which comes a different checker, `MallocChecker`) also gets suppressed.

```
void bad() {  
    auto smart = std::unique_ptr<int>(new int(13));  
    auto *raw = new int(29);  
    // There is a leak here.  
    // But that warning gets suppressed!  
}
```

But as a side-effect, the following code's leak warning (which comes a different checker, `MallocChecker`) also gets suppressed.

```
void bad() {  
    auto smart = std::unique_ptr<int>(new int(13));  
    auto *raw = new int(29);  
    // There is a leak here.  
    // But that warning gets suppressed!  
}
```

Moral of the story: Model C++ classes completely to avoid such side-effects.

# Closure

# Closure

- We have a almost complete model for `std::unique_ptr` in `SmartPointerModelling`



# Closure

- We have a almost complete model for `std::unique_ptr` in `SmartPointerModelling`
- The current modelling provides a framework for modelling other smart-pointers, such as `std::shared_ptr` or `std::weak_ptr`.

# Closure

- We have a almost complete model for `std::unique_ptr` in `SmartPtrModelling`
- The current modelling provides a framework for modelling other smart-pointers, such as `std::shared_ptr` or `std::weak_ptr`.
- New developers are very much welcome! Please feel free to hit us up on the cfe-dev mailing list.