



arm

Clang vs. GCC for SPEC2017 on AArch64

Sjoerd Meijer

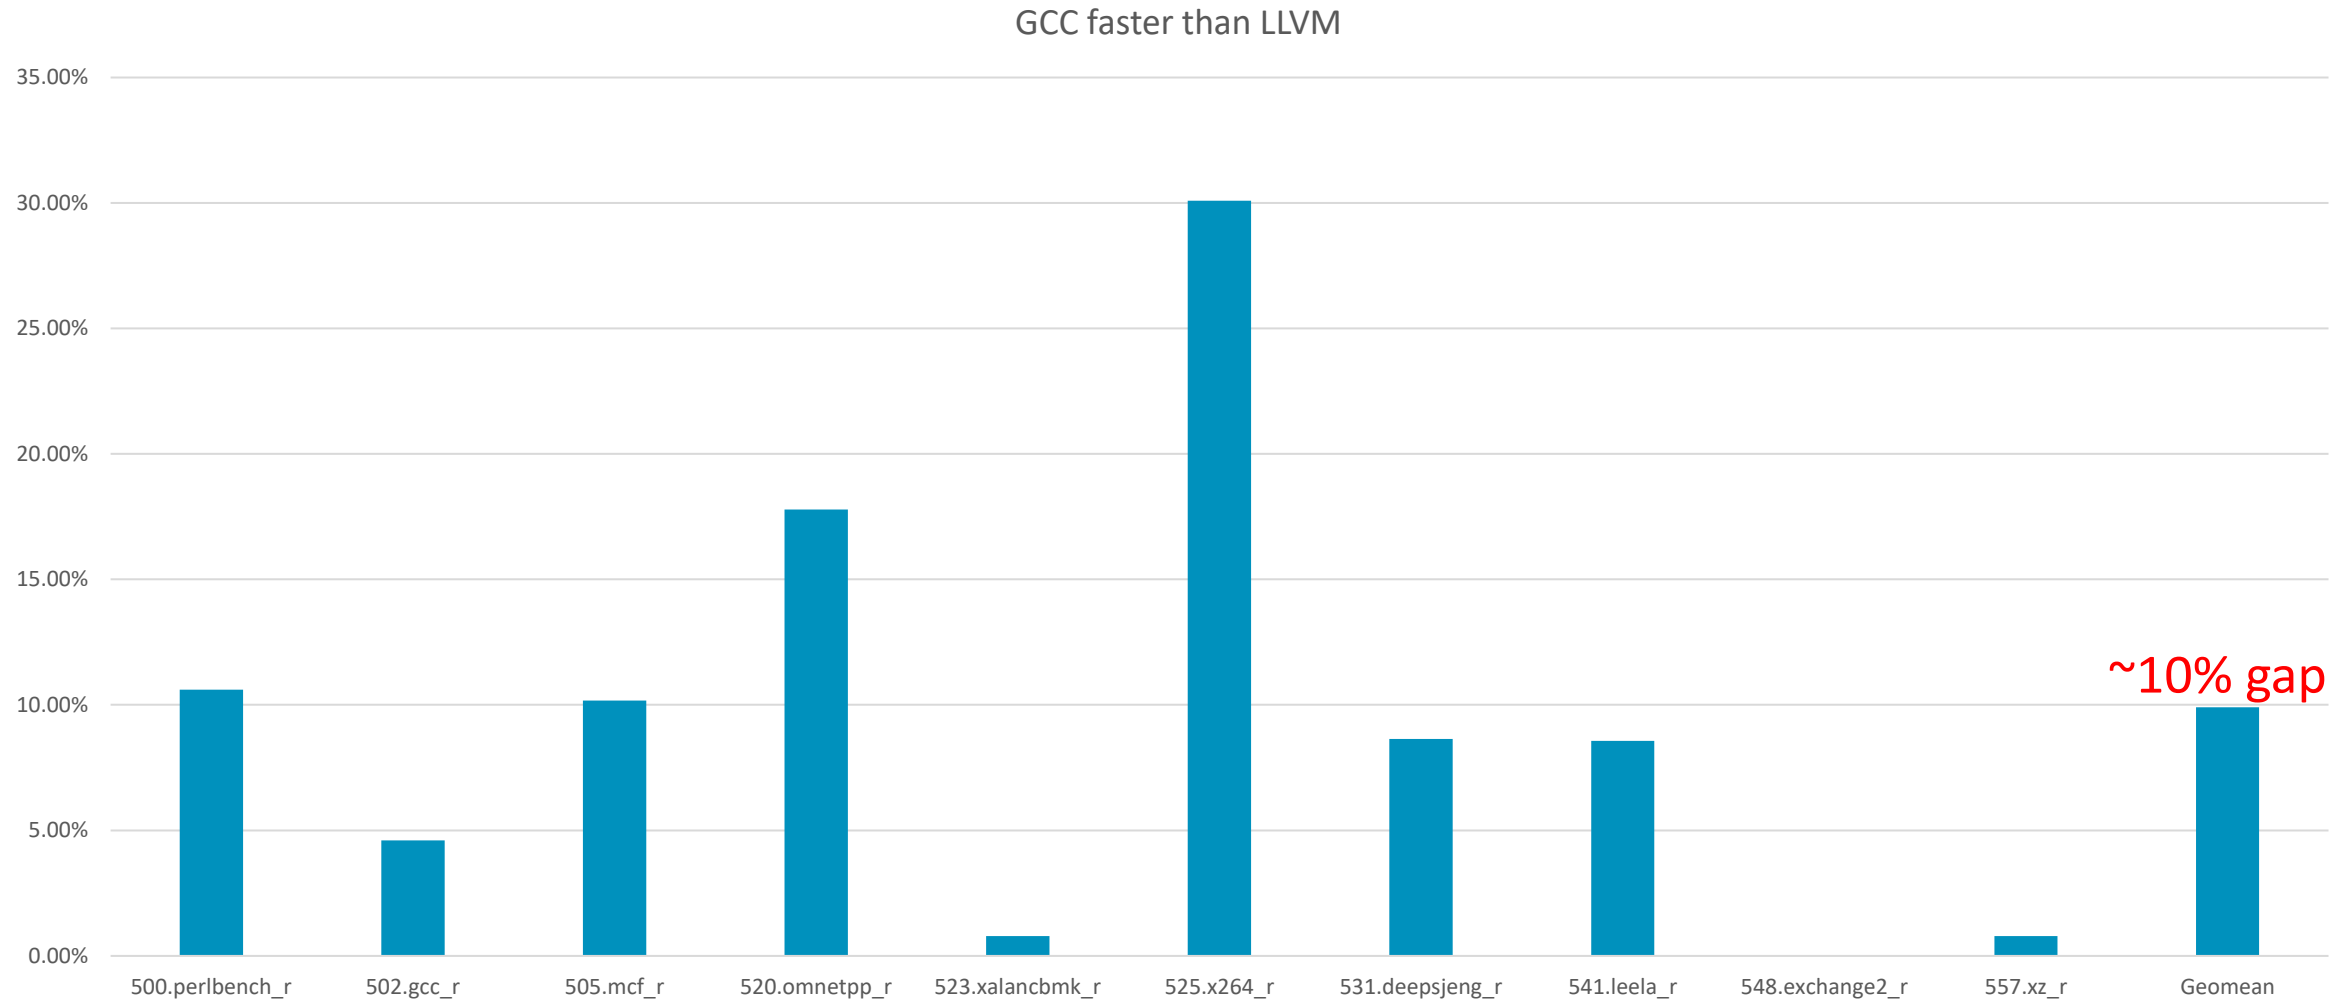
Sjoerd.meijer@arm.com

US LLVM Developer conference 2021


SPEC2017

- "performance measurements that can be used to compare compute-intensive workloads on different computer systems."
- Family of benchmarks: SPEC2006, SPEC2017.
 - Sub-suites for measuring integer and floating-point performance.
 - We will focus on SPEC2017Int rate.
- Contribute generic improvements, i.e. not only improve SPEC.
 - Target AArch64, but
 - will present work on target independent optimisations.

LLVM12 vs. GCC11 –Ofast




LLVM12 vs. GCC11 –Ofast



Fixed

- [D99354](#): Loopunswitching (omnetpp), JinGu.
 - Hoist loop-invariant control flow.
 - Port partially invariant unswitching from LoopUnswitch to SimpleLoopUnswitch.
 - Allows elimination of a dead loop.
- [D107281](#): Added a new pass function specialisation (MCF), Sjoerd.
 - Not yet enabled by default.
- Smaller auto-vec improvements (x264)
 - Added a match pattern, Dave
 - Loads of i8 vectors, and cost-modeling, Sjoerd


LLVM12 vs. GCC11 –Ofast, update



Fixed, cont'd

- [D107281](#): SimplifyCFG (leela), Momchil.
 - Speculate a store preceded by a local non-escaping load.
- Loopunrolling (perlbench), JinGu:
 - [D107381](#): Extends runtime unrolling for loops with multiple exiting blocks.
 - [D105996](#): Enable unrolling of non-constant bounds for all AArch64 targets.

LLVM12 vs. GCC11 –Ofast, update #2



Observations

- Target independent IR passes and optimisations
 - I.e., nothing AArch64 specific.
- Identified passes that:
 - Heavily interact with each other, e.g. unroller and the vectorisers.
 - Are missing functionality, e.g. the SLP vectoriser.
 - Are not enabled by default, e.g. GVNHoist, function specialisation.
 - Are not generic enough, e.g. loop-distribute.

Pass Ordering Challenges

- X264 has kernels with small, constant integer bounds:
 - Fully unrolled first -> Loop vectoriser doesn't run -> SLP vectoriser is missing features.

```
void fn( int16_t *A, uint16_t *B, uint16_t *C )
{
    for( int i = 0; i < 16; i++ ) {
        if (A[i] > 0 )
            A[i] = (C[i] + A[i]) * B[i] >> 16;
        else
            A[i] = - ((C[i] - A[i]) * B[i] >> 16);
    }
}
```

- [D102748](#): Don't unroll before vectorisation.
 - PR47178, PR47726, PR47554, PR47436, PR31572, PR47553, PR47491.
 - Consensus:
 - full unrolling happens early: scalar optimizations have a chance to work on fully unrolled loops
 - SLP vectorizer needs fixing.

Passes Lacking Features

- [SLPVectorizer] Implement initial memory versioning: [D102834](#)
 - Emit runtime alias checks, like the loop vectoriser.

```
void fn( int16_t *A, uint16_t *B, uint16_t *C )
{
    for( int i = 0; i < 16; i++ ) {
        if (A[i] > 0 )
            A[i] = (C[i] + A[i]) * B[i] >> 16;
        else
            A[i] = - ((C[i] - A[i]) * B[i] >> 16);
    }
}
```

- Other, target dependent optimisations:
 - If-conversion
 - Cost-modeling

Passes Not Enabled

- Function Specialisation:
 - TODO: Is it generic enough, is worth the extra compile-times?
- GVNHoist:
 - Enabled/reverted a few times, still correctness and perf problems
 - [GVN] Simple GVN hoist – scalars: [D110817](#)
 - [GVN] Simple GVN hoist - loads and stores: [D110822](#)

Passes Not Generic Enough

- Loop-distribute:
 - Helps hammer in SPEC2006.
 - Improve loop distribute cost model: [D100381](#)
 - But were not able to make this generic.

Pass Missing?

- Loop Invariant Code Motion (LICM)
 - Philosophy: canonicalisation transformation.
 - Aggressively hoists everything it can, backend should undo this.
 - Except, we don't have a good solution:
 - LoopSink: works on IR and only with profile info
 - MachineSink: can't sink into loops.

Various other TODOs

- Missed opportunities for register promotion: [PR51193](#)
- Multiple evaluations of a GEP: [PR51184](#)
- Sub-optimal placement of loop exit blocks: [PR51185](#)
- Failure to recognize table-based ctz implementation: [PR46434](#)

Conclusion

- Have some way to go:
 - Auto-vectorisation (SLP vectoriser)
 - Hoisting
 - Undoing LICM?
 - Rinse & Repeat the whole process
- Feedback welcome!
 - Please get in touch on the dev list, phabricator, or email.

arm

Thank You

Danke

Merci

谢谢

ありがとう

Gracias

Kiitos

감사합니다

धन्यवाद

شكراً

תודה