

OPENCL SUPPORT IN CLANG STATUS : OPENCL C 3.0, IMPROVEMENTS AND FUTURE DIRECTIONS

Anton Zabaznov

OPENCL C 3.0 AND OPTIONAL FUNCTIONALITY

- OpenCL is a standard for heterogeneous computing.
- [OpenCL C 3.0 specification](#) was released on September 30th 2020.
- OpenCL C 3.0 is backward compatible with 1.2 and introduces optionality of 2.0 features (such as device enqueue, program scope variables etc).
- This brings flexibility for vendors to implement necessary functionality targeting developer requests.



OPENCL C 3.0 AND OPTIONAL FUNCTIONALITY

- Support of various optional features in OpenCL C 3.0 is indicated with a feature test macros within a kernel source ([full list of optional features](#)):

Feature macro	Description
<code>__opencl_c_3d_image_writes</code>	Supports of 3D image type and built-in functions for writing to 3D image objects.
<code>__opencl_c_device_enqueue</code>	Supports of necessary types and built-in functions to enqueue additional work from the device.
<code>__opencl_c_generic_address_space</code>	Support of the unnamed generic address space.
<code>__opencl_c_program_scope_global_variables</code>	Support of program scope variables in the global address space.
...	...

OPENCL C 3.0 SUPPORT IN CLANG

- There already exists such concept in OpenCL as *extensions* (see [OpenCL extension specification](#)). Handling of OpenCL *extensions* in clang was highly reused to implement OpenCL C 3.0 support.
 - Mostly all features are implemented in clang (see [OpenCL C 3.0 Implementation Status](#)).
 - Existing features are set by the target or can be overridden via existing command line flag `-cl-ext`:
- ```
$ clang -cl-std=CL3.0 -target r600 -Xclang -cl-ext=__opencl_c_generic_address_space ...
```
- Some work is still in progress regarding support of device enqueue and adding of built-in functions.

# DEVICE ENQUEUE WITHOUT PROGRAM SCOPE VARIABLES FEATURE

- Device enqueue syntax:

```
kernel void kernel foo(uint local_mem_size) {
 void (^block_var)(local void*) // block variable
 = ^void(void local *ptr) {...}; // block literal expression
 enqueue_kernel(..., block_var, ..., local_mem_size) // builtin function
}
```

- Block variable is initialized with block literal expression.
- Block literal expression is a struct, for OpenCL it contains a pointer to function for invocation and some other fields.

# DEVICE ENQUEUE WITHOUT PROGRAM SCOPE VARIABLES FEATURE

- Block literal expression for block with no captures is emitted in global scope as optimization according to blocks ABI:

```
$ clang -cl-std=CL3.0 -Xclang
-cl-ext=__openc1_c_device_enqueue,__openc1_c_program_scope_global_variables
```

```
void foo() {
 ...
 int (^b)(int) = ^(int n) {
 return n * n;
 };
 ...
}
```

```
__openc1_block_literal_generic = type {i32, i32, i8 addrspace(4)*}

@__block_literal_global = internal addrspace(1)
 constant {i32, i32, i8 addrspace(4)*} ...

spir_func void @foo() {
 %1 = alloca __openc1_block_literal_generic addrspace(4)*
 store (cast @__block_literal_global
 to %__openc1_block_literal_generic addrspace(4)*),
 __openc1_block_literal_generic addrspace(4)** %1
 ...
}
```

- Now global address space is used. That's valid if feature for program scope global variables is supported.

# DEVICE ENQUEUE WITHOUT PROGRAM SCOPE VARIABLES FEATURE

- Otherwise, block literal expression is emitted in local scope (in progress, [D112230](#)).

```
$ clang -cl-std=CL3.0 -Xclang
-cl-ext=__openc1_c_device_enqueue,-__openc1_c_program_scope_global_variables
```

```
void foo() {
 ...
 int (^b)(int) = ^(int n) {
 return n * n;
 };
 ...
}
```

```
__openc1_block_literal_generic = type {i32, i32, i8 addrspace(4)*}

spir_func void @foo() {
 %1 = alloca __openc1_block_literal_generic addrspace(4)*
 %2 = alloca <{ i32, i32, i8 addrspace(4)* }>
 ; local block literal initialization
 %6 = cast %2 to __openc1_block_literal_generic addrspace(4)*
 store __openc1_block_literal_generic addrspace(4)* %6,
 __openc1_block_literal_generic addrspace(4)** %1
 ...
}
```

- Some refactoring of code generation may be needed for global blocks to support constant address space.

# REDUNDANT AST NODES FOR IMPLICIT DEFINITIONS

```
$ clang -cl-std=CL2.0 -Xclang -ast-dump
```

```
kernel void foo() {
 ...
};
```



AST

```
-TypedefDecl 0xcd180 implicit atomic_long '_Atomic(long)'
 |-AtomicType 0xcd130 '_Atomic(long)'
-TypedefDecl 0xcd230 implicit atomic_ulong '_Atomic(unsigned long)'
 |-AtomicType 0xcd1e0 '_Atomic(unsigned long)'
-TypedefDecl 0xcd390 implicit atomic_double '_Atomic(double)'
 |-AtomicType 0xcd340 '_Atomic(double)'
-TypedefDecl 0xcd570 implicit atomic_size_t '_Atomic(unsigned long)'
 |-AtomicType 0xcd1e0 '_Atomic(unsigned long)'
...
```



# REDUNDANT AST NODES FOR IMPLICIT DEFINITIONS

\$ clang -cl-std=CL2.0 -Xclang -ast-dump

```
kernel void foo() {
 ...
};
```



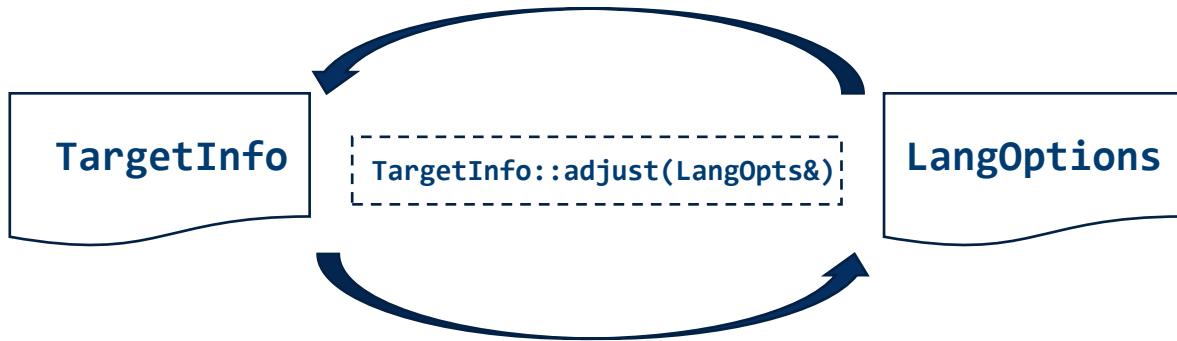
AST

```
-TypedefDecl 0xcd180 implicit atomic_long '_Atomic(long)'
 ^_AtomicType 0xcd130 '_Atomic(long)'
-TypedefDecl 0xcd230 implicit atomic_ulong '_Atomic(unsigned long)'
 ^_AtomicType 0xcd1e0 '_Atomic(unsigned long)'
-TypedefDecl 0xcd390 implicit atomic_double '_Atomic(double)'
 ^_AtomicType 0xcd340 '_Atomic(double)'
-TypedefDecl 0xcd570 implicit atomic_size_t '_Atomic(unsigned long)'
 ^_AtomicType 0xcd1e0 '_Atomic(unsigned long)'
...
```

- Removed AST nodes creation for implicit definitions when unnecessary and simplified diagnostics: D97058, D92244

# TARGET OPTIONS AND LANGUAGE OPTIONS SETTING IN CLANG

- TargetInfo and LangOptions augment each other:



- Some target options are set according to language options.
- Some language options are set according to target information, for example arithmetic fences, OpenCL generic address space/pipes, AltiVec extensions in PPC target.

# TARGET OPTIONS AND LANGUAGE OPTIONS SETTING IN CLANG

- As a result, TargetInfo is not immutable after creation:

CreateTargetInfo(TargetOptions&)

TargetInfo

```
StringMap<bool> FeatureMap;
```

```
...
```

```
virtual void adjust(LangOpts&);
```

```
...
```



CompilerInstance::createTarget()

```
...
```

```
TI = CreateTargetInfo(TargetOpts)
```

```
...
```

```
TI.adjust(LangOpts);
```



```
...
```

- TargetInfo initialized via TargetOptions

- LangOpts are changed with TargetInfo
- TargetInfo is changed with LangOpts

**Target is not immutable after creation as it meant to be!**

# TARGET OPTIONS AND LANGUAGE OPTIONS SETTING IN CLANG

```
CreateTargetInfo(TargetOptions&,
 const LangOpts&)
```

## TargetInfo

```
StringMap<bool> FeatureMap;

...
virtual void adjustWithLO
 (const LangOpts&);
...
```



```
CompilerInstance::createTarget()
```

```
...
TI = CreateTargetInfo(TargetOpts,
 LangOpts)
...
LangOpts.adjustWithTI(TI);
```



...

- TargetInfo initialized via:
  - TargetOptions
  - LangOptions

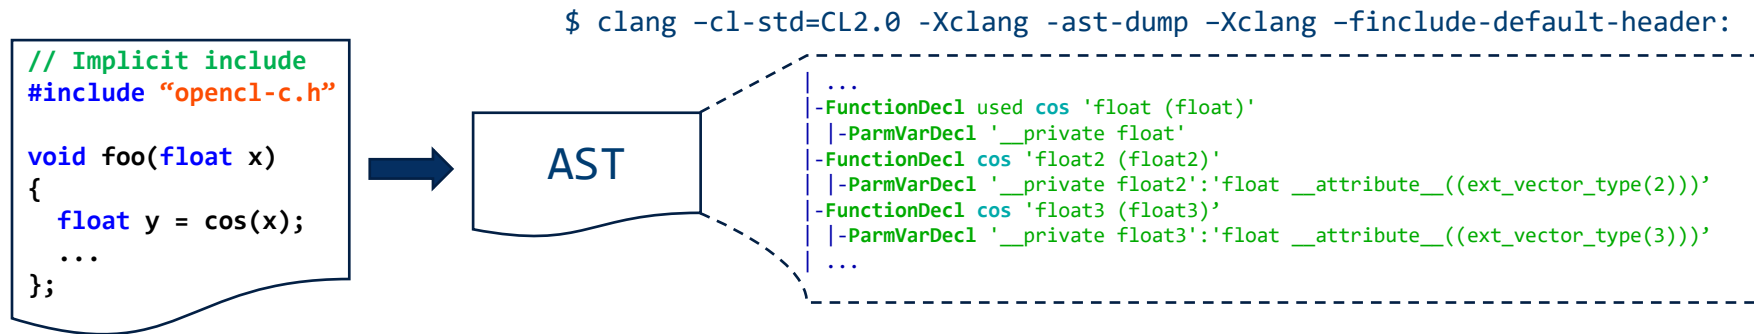
- LangOpts are changed with TargetInfo

Target is immutable after creation

- Proposed design solution to extend interfaces when creating a target info: [D110036](#)

# OTHER IMPROVEMENTS AND FUTURE DIRECTIONS: BUILT-INS


- Old behaviour:



- To use OpenCL built-in functions, user had to explicitly specify `-finclude-default-header` flag, which includes a file with all OpenCL built-in functions declarations.
- As a result, increase in compilation time when parsing header with declarations.

# OTHER IMPROVEMENTS AND FUTURE DIRECTIONS: BUILT-INS

- New behaviour:



- Dynamically insert built-in function declaration from table if needed. Table is auto-generated from TableGen file.
- Implemented by Sven Van Haastregt, see [his talk](#).

# OTHER IMPROVEMENTS AND FUTURE DIRECTIONS

- Dynamic lookup is enabled as default behaviour: [D96515](#).
  - Option `-cl-no-stdinc` can be used to use the old approach (via header).
- Planning to deprecate usage of `opencl-c.h` and always use dynamic lookup.
  - Testing scheme for soft switching is proposed: [D99577](#)
  - Make sure that TableGen definitions provide on-to-one functionality with `opencl-c.h` via diffing.
- Planning to support device enqueue for OpenCL C 3.0.
- Planning to support the rest of OpenCL C 3.0 built-ins.

# CONTRIBUTORS

- Anastasia Stulova, Arm
- Sven van Haastregt, Arm
- Anton Zabaznov, Intel
- Dave Airlie, Red Hat