

# Optimizing OpenMP GPU Execution in LLVM

Melanie Cornelius (IIT), Jose Monsalve Diaz (ANL), Johannes Doerfert (ANL),  
Giorgis Georgakoudis (LLNL), Joseph Huber (ORNL),  
Atmn Patel (U of Waterloo), Shilei Tian (Stony Brook University)

Background

# OpenMP in LLVM

Weekly Meeting: <https://bit.ly/2Zqt49v>

## Flang

### Clang

OpenMP  
Parser

OpenMP  
Sema

OpenMP  
CodeGen

### OpenMPIRBuilder

frontend-independent  
OpenMP LLVM-IR  
generation

favor simple and  
expressive LLVM-IR

reusable for non-OpenMP  
parallelism

### OpenMPOpt

interprocedural  
optimization pass

contains host & device  
optimizations

run with -O1 and -O2  
since LLVM 11



### OpenMP runtimes

libomp.so (host)

libomptarget + plugins

(offloading, host)

libomptarget-nvptx  
(offloading, device)

# OpenMP in LLVM

Weekly Meeting: <https://bit.ly/2Zqt49v>

Flang

Clang

OpenMP  
Parser

OpenMP  
Sema

OpenMP  
CodeGen

OpenMPIRBuilder

frontend-independent  
OpenMP LLVM-IR  
generation

favor simple and  
expressive LLVM-IR

reusable for non-OpenMP  
parallelism

New  
optimizations  
for GPUs

OpenMPOpt

interprocedural  
optimization pass

contains host & **device  
optimizations**

run with -O1 and -O2  
since LLVM 11

OpenMP  
runtimes

libomp.so (host)

libomptarget + plugins

(offloading, host)

libomptarget-nvptx  
(offloading, device)



# OpenMP in LLVM

Weekly Meeting: <https://bit.ly/2Zqt49v>

Flang

Clang

OpenMP  
Parser

OpenMP  
Sema

OpenMP  
CodeGen

OpenMPIRBuilder

frontend-independent  
OpenMP LLVM-IR  
generation

favor simple and  
expressive LLVM-IR

reusable for non-OpenMP  
parallelism

New  
optimizations  
for GPUs

OpenMPOpt

interprocedural  
optimization pass

contains host & device  
optimizations

run with -O1 and -O2  
since LLVM 11

Re-design for  
optimization

OpenMP  
runtimes

libomp.so (host)

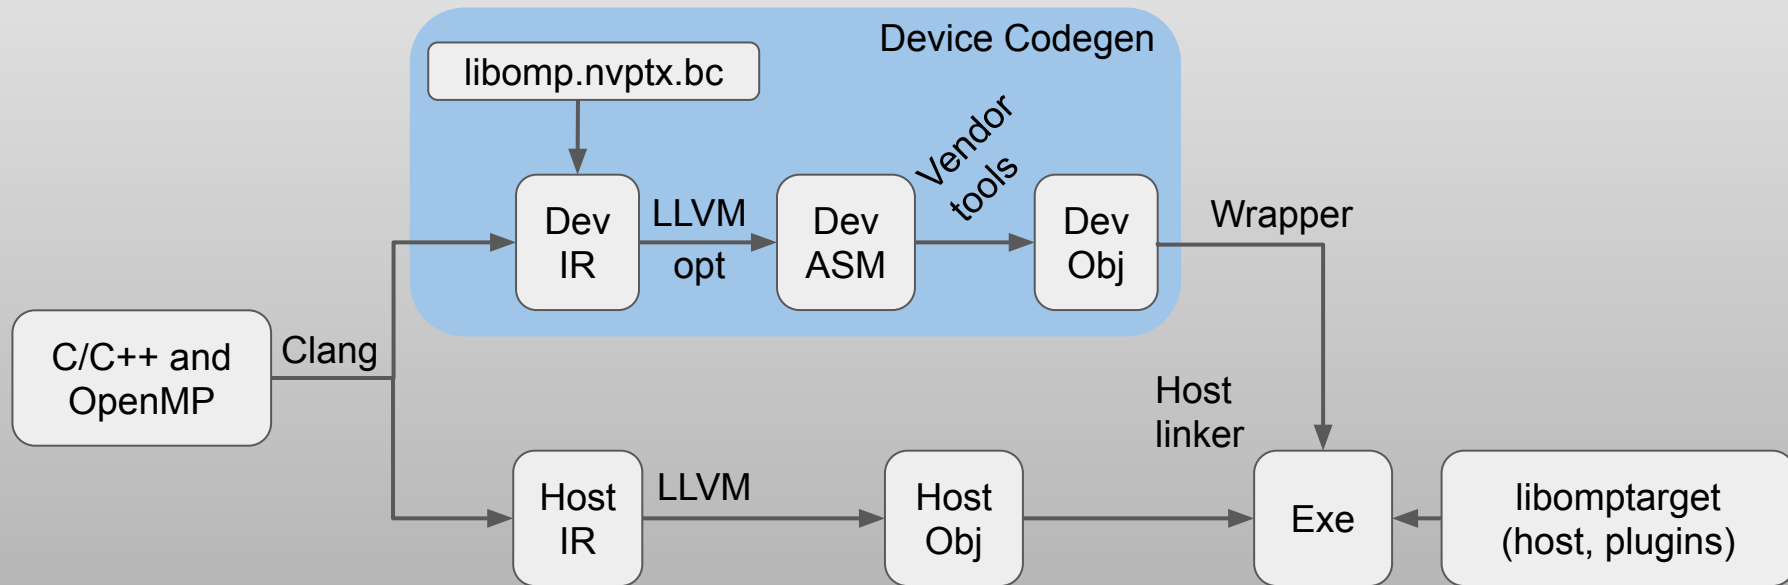
libomptarget + plugins

(offloading, host)

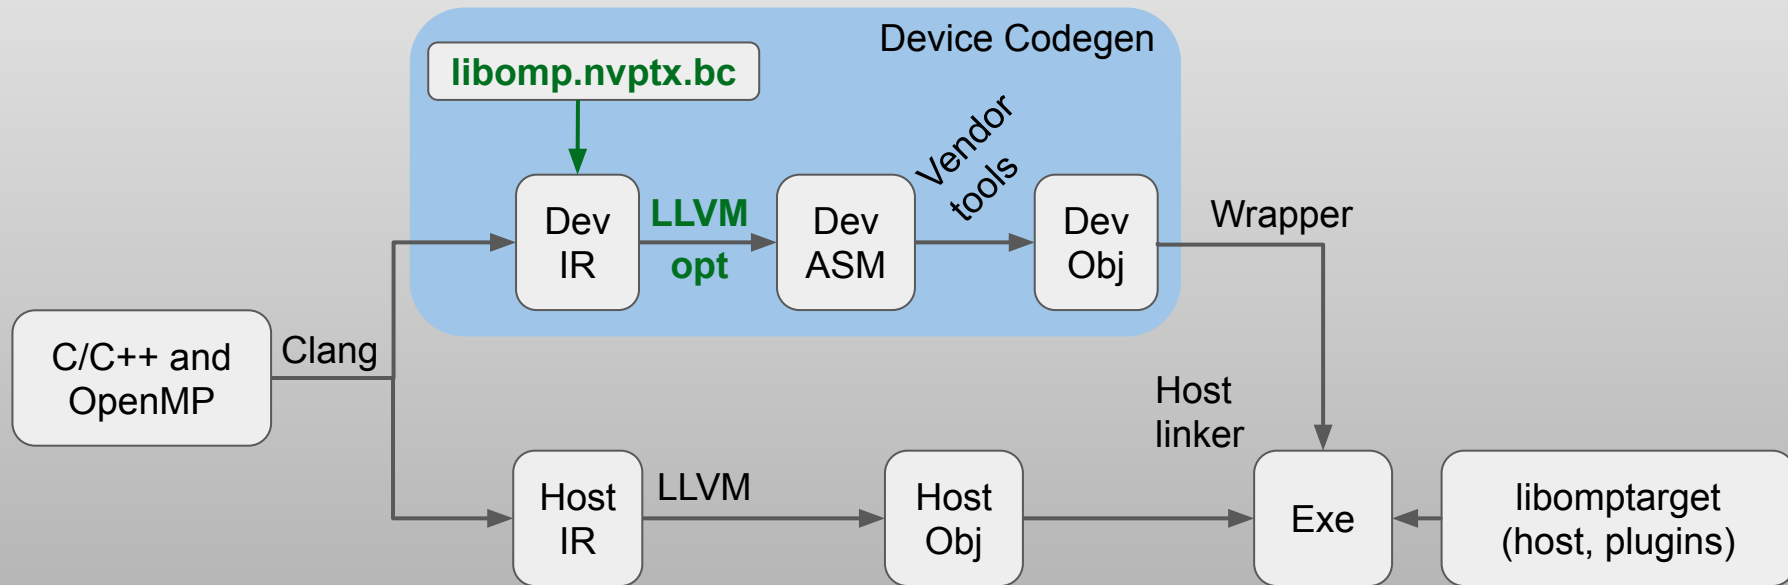
libomptarget-nvptx  
(offloading, device)




# OpenMP Offload Compilation Toolchain



# OpenMP Offload Compilation Toolchain



# OpenMP to GPU Mapping





# (Some) Motivational Problems

# OpenMP Offload vs CUDA

```
__global__ void cuda() {  
  
    __shared__ double Buffer[BLOCK_SIZE];  
  
    int L;  
  
    if (threadIdx.x == 0)  
        single_thread_init();  
    __syncthreads();  
  
    L = load_data(Buffer, threadIdx.x);  
    __syncthreads();  
  
    if (L != 0)  
        parallel_work(Buffer, threadIdx.x);  
  
}
```

```
void openmp() {  
    #pragma omp target teams distribute  
    for (int i = 0; i < GRID_SIZE; i += BLOCK_SIZE) {  
        double Buffer[BLOCK_SIZE];  
  
        int L;  
  
        // No conditional, conceptually one thread only  
        single_thread_init();  
        // No synchronization, again, one thread  
  
        #pragma omp parallel for  
        for (int j = 0; j < BLOCK_SIZE; ++j)  
            L = load_data(Buffer, j);  
        // Synchronization is implicit  
  
        #pragma omp parallel for  
        for (int j = 0; j < BLOCK_SIZE; ++j)  
            if (L != 0)  
                parallel_work(Buffer, j);  
  
    }  
}
```

# OpenMP Offload vs CUDA - Globalization of Locals

```
__global__ void cuda() {
```

```
    __shared__ double Buffer[BLOCK_SIZE];
```

```
    int L;
```

```
    if (threadIdx.x == 0)
        single_thread_init();
    __syncthreads();
```

```
    L = load_data(Buffer, threadIdx.x);
    __syncthreads();
```

```
    if (L != 0)
        parallel_work(Buffer, threadIdx.x);
```

```
}
```



```
void openmp_impl() {
    #pragma omp target teams distribute
    for (int i = 0; i < GRID_SIZE; i += BLOCK_SIZE) {
        double *Buffer = __omp_alloc(8 * BLOCK_SIZE);
```

```
        int *L = __omp_alloc(sizeof(int));
```

```
        // No conditional, conceptually one thread only
        single_thread_init();
        // No synchronization, again, one thread
```

```
        #pragma omp parallel for shared(L)
        for (int j = 0; j < BLOCK_SIZE; ++j)
            *L = load_data(Buffer, j);
        // Synchronization is implicit
```

```
        #pragma omp parallel for shared(L)
        for (int j = 0; j < BLOCK_SIZE; ++j)
            if (*L != 0)
                parallel_work(Buffer, j);
```

```
        // __omp_free(...)
    }
```


```
}
```



# OpenMP Offload vs CUDA - Execution Mode

```
__global__ void cuda() {
```


## SPMD/GPU Execution Mode



```
if (L != 0)
    parallel_work(Buffer, threadIdx.x);
```

```
void openmp_impl() {
#pragma omp target teams distribute
for (int i = 0; i < GRID_SIZE; i += BLOCK_SIZE) {
```

## Generic/CPU Execution Mode



```
    pragma omp parallel for
    for (int j = 0; j < BLOCK_SIZE; ++j)
        if (*L != 0)
            parallel_work(Buffer, j);
```

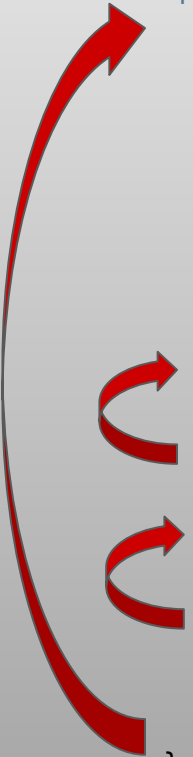
```
    // __omp_free(...)
```

```
    }
}
```

# OpenMP Offload vs CUDA - Explicit Loop Backedges

```
__global__ void cuda() {  
  
    __shared__ double Buffer[BLOCK_SIZE];  
  
    int L;  
  
    if (threadIdx.x == 0)  
        single_thread_init();  
    __syncthreads();  
  
    L = load_data(Buffer, threadIdx.x);  
    __syncthreads();  
  
    if (L != 0)  
        parallel_work(Buffer, threadIdx.x);  
  
}
```

```
void openmp_impl() {  
    #pragma omp target teams distribute  
    for (int i = 0; i < GRID_SIZE; i += BLOCK_SIZE) {  
        double *Buffer = __omp_alloc(8 * BLOCK_SIZE);  
  
        int *L = __omp_alloc(sizeof(int));  
  
        // No conditional, conceptually one thread only  
        single_thread_init();  
        // No synchronization, again, one thread  
  
        #pragma omp parallel for  
        for (int j = 0; j < BLOCK_SIZE; ++j)  
            *L = load_data(Buffer, j);  
        // Synchronization is implicit  
  
        #pragma omp parallel for  
        for (int j = 0; j < BLOCK_SIZE; ++j)  
            if (*L != 0)  
                parallel_work(Buffer, j);  
  
        // __omp_free(...)  
    }  
}
```



# Optimizations

# OpenMP Offloading - Deglobalization

```
__global__ void cuda() {  
  
    __shared__ double Buffer[BLOCK_SIZE];  
  
    int L;  
  
    if (threadIdx.x == 0)  
        single_thread_init();  
    __syncthreads();  
  
    L = load_data(Buffer, threadIdx.x);  
    __syncthreads();  
  
    if (L != 0)  
        parallel_work(Buffer, threadIdx.x);  
  
}
```

```
void openmp_impl() {  
    #pragma omp target teams distribute  
  
    for (int i = 0; i < GRID_SIZE; i += BLOCK_SIZE) {  
        double Buffer[BLOCK_SIZE];  
        #pragma omp allocate(Buffer) allocator(cgroup)  
        int L;  
        #pragma omp allocate(L) allocator(thread)  
        // No conditional, conceptually one thread only  
        single_thread_init();  
        // No synchronization, again, one thread  
  
        #pragma omp parallel for  
        for (int j = 0; j < BLOCK_SIZE; ++j)  
            L = load_data(Buffer, j);  
        // Synchronization is implicit  
  
        #pragma omp parallel for  
        for (int j = 0; j < BLOCK_SIZE; ++j)  
            if (L != 0)  
                parallel_work(Buffer, j);  
    }  
}
```

# OpenMP Offloading - SPMDzation

```
__global__ void cuda() {  
  
    __shared__ double Buffer[BLOCK_SIZE];  
  
    int L;  
  
    if (threadIdx.x == 0)  
        single_thread_init();  
    __syncthreads();  
  
    L = load_data(Buffer, threadIdx.x);  
    __syncthreads();  
  
    if (L != 0)  
        parallel_work(Buffer, threadIdx.x);  
  
}
```

```
void openmp_impl() {  
    #pragma omp target teams distribute  
    #pragma omp parallel  
        for (int i = 0; i < GRID_SIZE; i += BLOCK_SIZE) {  
            double Buffer[BLOCK_SIZE];  
            #pragma omp allocate(Buffer) allocator(cgroup)  
            int L;  
            #pragma omp allocate(L) allocator(thread)  
            if (__omp_get_thread_id() == 0)  
                single_thread_init();  
            #pragma omp barrier // aligned  
  
            #pragma omp for nowait  
            for (int j = 0; j < BLOCK_SIZE; ++j)  
                L = load_data(Buffer, j);  
            #pragma omp barrier // aligned  
  
            #pragma omp for nowait  
            for (int j = 0; j < BLOCK_SIZE; ++j)  
                if (L != 0)  
                    parallel_work(Buffer, j);  
            #pragma omp barrier // aligned  
        }  
}
```



# OpenMP Offloading - Loop Oversubscription

```
__global__ void cuda() {  
  
    __shared__ double Buffer[BLOCK_SIZE];  
  
    int L;  
  
    if (threadIdx.x == 0)  
        single_thread_init();  
    __syncthreads();  
  
    L = load_data(Buffer, threadIdx.x);  
    __syncthreads();  
  
    if (L != 0)  
        parallel_work(Buffer, threadIdx.x);  
  
}
```

```
void openmp_impl() {  
#pragma omp target teams parallel  
int i = omp_get_team_num();  
if (i < GRID_SIZE) { // known true  
    double Buffer[BLOCK_SIZE];  
    #pragma omp allocate(Buffer) allocator(cgroup)  
    int L;  
    #pragma omp allocate(L) allocator(thread)  
    if (__omp_get_thread_id() == 0)  
        single_thread_init();  
    #pragma omp barrier // aligned  
  
    int j = omp_get_thread_num();  
    if (j < BLOCK_SIZE) // known true  
        L = load_data(Buffer, j);  
    #pragma omp barrier // aligned  
  
    int j = omp_get_thread_num();  
    if (j < BLOCK_SIZE) // known true  
        if (L != 0)  
            parallel_work(Buffer, j);  
    #pragma omp barrier // aligned  
}  
}
```

# OpenMP Offloading - (Aligned) Barrier Removal

```
__global__ void cuda() {  
  
    __shared__ double Buffer[BLOCK_SIZE];  
  
    int L;  
  
    if (threadIdx.x == 0)  
        single_thread_init();  
    __syncthreads();  
  
    L = load_data(Buffer, threadIdx.x);  
    __syncthreads();  
  
    if (L != 0)  
        parallel_work(Buffer, threadIdx.x);  
  
}
```

```
void openmp_impl() {  
    #pragma omp target teams parallel  
    int i = omp_get_team_num();  
    if (i < GRID_SIZE) { // known true  
        double Buffer[BLOCK_SIZE];  
        #pragma omp allocate(Buffer) allocator(cgroup)  
        int L;  
        #pragma omp allocate(L) allocator(thread)  
        if (__omp_get_thread_id() == 0)  
            single_thread_init();  
        #pragma omp barrier // aligned  
  
        int j = omp_get_thread_num();  
        if (j < BLOCK_SIZE) // known true  
            L = load_data(Buffer, j);  
        #pragma omp barrier // aligned  
  
        int j = omp_get_thread_num();  
        if (j < BLOCK_SIZE) // known true  
            if (L != 0)  
                parallel_work(Buffer, j);  
        #pragma omp barrier // aligned  
    }  
}
```

# Optimization Implementation

# The OpenMP-Opt Pass


- OpenMP-specific optimizations (= instant no-op for non-openmp codes)
- Run early (as module pass) and late (as CG-SCC pass)
- Embedded Domain Knowledge (recognizes `omp_*` and `__kmpc_*` calls)
- Uses the Attributor IPO framework and provides custom Abstract Attributes:
  - **AAExecutionDomain** - Determine if a block is executed by the main thread only, or by all threads in an “aligned” fashion.
  - **AAFoldRuntimeCall** - Replace runtime calls with their constant return value (if known).
  - **AAHeapToStack** - Replace globalized memory with an `alloca`.
  - **AAHeapToShared** - Replace globalized memory with shared memory (if heap-2-stack failed).
  - **AAKernellInfo** - Track reaching kernels, optimize kernel execution mode, ...
- *Inter-procedural by design*

# Inter-procedural analysis with OpenMP-awareness

- Internalize functions to improve inter-procedural analysis
  - All calls to the internalized version are known
- Analyse uses of known OpenMP runtime calls

```
define void @__omp_offloading_XXX() {  
  call void @foo()  
  ret void  
}
```


Replace call



```
define void @__omp_offloading_XXX() {  
  call void @foo.internalized()  
  ret void  
}
```

```
define void @foo() {  
entry:  
  %0 = call i8* @__omp_alloc(i64 4)  
  call void @bar(i8* %0)  
  call void @__omp_free(i8* %0)  
  ret void  
}
```

Clone function



```
define internal void @foo.internalized() {  
entry:  
  %0 = call i8* @__omp_alloc(i64 4)  
  call void @bar.internalized(i8* %0)  
  call void @__omp_free(i8* %0)  
  ret void  
}
```

Analyze calls,  
replace uses



```
{  
  %0 = call i8* @__omp_alloc(i64 4)  
  call void @bar.internalized(i8* %0)  
  call void @__omp_free(i8* %0)  
  ret void  
}
```

# Runtime Co-Design

# Remarks and Assumptions for Interactive Optimization

OpenMP-Opt emits **remarks**:

- ❑ -Rpass=openmp-opt
- ❑ -Rpass-missed=openmp-opt
- ❑ -Rpass-analysis=openmp-opt

to report success and failure,  
and utilizes **assumptions**:

- ❑ `#pragma omp assumes ...`
- ❑ `__attribute__((assume("...")))`
- ❑ command line flags

to enhance static analysis.

`omp_no_openmp`

`omp_no_parallelism`

`omp_no_openmp_routines`

} OpenMP 5.1 spec  
assumptions

`ompx_spm_d_amenable`

`ompx_aligned_barrier`

`ompx_no_sync`

} LLVM assumption  
extensions

`-fopenmp-assume-teams-oversubscription`

`-fopenmp-assume-threads-oversubscription`

# Explicit (Shared) Global State and Powerful IPO

```
__shared__ StateTy State;
```

```
__global__ void kernel() {  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_parallel(outlined_fn, ...);  
}
```

```
__device__ static void outlined_fn(...) {  
    // Do not (transitively) call __omp_parallel.  
    use(State.TeamSize);  
}
```



# Explicit (Shared) Global State and Powerful IPO

```
__shared__ StateTy State;

__global__ void kernel() {
    State.TeamSize = 1;
    __omp_aligned_barrier(); // assume((ompx_aligned))
    __omp_parallel(outlined_fn, ...);
}

__device__ static void __omp_parallel(fn, ...) {
    if (State.TeamSize > 1)
        return __omp_parallel_sequentialized(fn, ...);
    __omp_aligned_barrier(); // assume((ompx_aligned))
    State.TeamSize = blockDim.x;
    __omp_aligned_barrier(); // assume((ompx_aligned))
    fn();
    __omp_aligned_barrier(); // assume((ompx_aligned))
    State.TeamSize = 1;
    __omp_aligned_barrier(); // assume((ompx_aligned))
}

__device__ static void outlined_fn(...) {
    // Do not (transitively) call __omp_parallel.
    use(State.TeamSize);
}
```

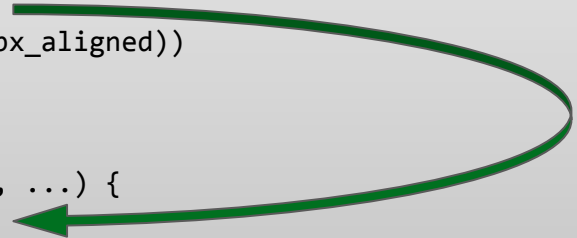
# Explicit (Shared) Global State and Powerful IPO

```
__shared__ StateTy State;
```

```
__global__ void kernel() {  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_parallel(outlined_fn, ...);  
}
```

```
__device__ static void __omp_parallel(fn, ...) {  
    if (1 > 1)  
        return __omp_parallel_sequentialized(fn, ...);  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    State.TeamSize = blockDim.x;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    fn();  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
}
```

```
__device__ static void outlined_fn(...) {  
    // Do not (transitively) call __omp_parallel.  
    use(State.TeamSize);  
}
```



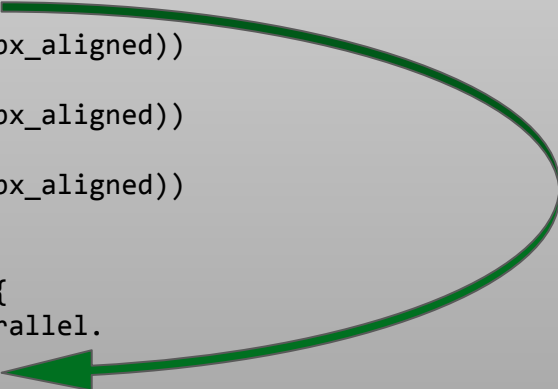
IP-Reachability +  
shared memory lifetime

# Explicit (Shared) Global State and Powerful IPO

```
__shared__ StateTy State;  
  
__global__ void kernel() {  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_parallel(outlined_fn, ...);  
}
```

```
__device__ static void __omp_parallel(fn, ...) {  
  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    State.TeamSize = blockDim.x;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    fn();  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
}
```

```
__device__ static void outlined_fn(...) {  
    // Do not (transitively) call __omp_parallel.  
    use(blockDim.x);  
}
```



IP-Reachability +  
shared memory lifetime +  
IP-Dominance +  
intrinsic annotations

# Explicit (Shared) Global State and Powerful IPO

```
__shared__ StateTy State;
```

← shared memory lifetime + IP-write-only

```
__global__ void kernel() {  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_parallel(outlined_fn, ...);  
}
```

←

```
__device__ static void __omp_parallel(fn, ...) {
```

```
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    State.TeamSize = blockDim.x;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    fn();  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    State.TeamSize = 1;  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
}
```

←

←

} shared memory lifetime + IP-DSE

```
__device__ static void outlined_fn(...) {  
    // Do not (transitively) call __omp_parallel.  
    use(blockDim.x);  
}
```

# Explicit (Shared) Global State and Powerful IPO

```
__global__ void kernel() {  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_parallel(outlined_fn, ...);  
}
```

```
__device__ static void __omp_parallel(fn, ...) {
```

```
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    fn();  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
    __omp_aligned_barrier(); // assume((ompx_aligned))  
}
```

```
__device__ static void outlined_fn(...) {  
    // Do not (transitively) call __omp_parallel.  
    use(blockDim.x);  
}
```



IP-aligned barrier elimination

# Explicit (Shared) Global State and Powerful IPO

```
__global__ void kernel() {  
  
    __omp_parallel(outlined_fn, ...);  
}  
  
__device__ static void __omp_parallel(fn, ...) {  
  
    fn();  
  
}  
  
__device__ static void outlined_fn(...) {  
    // Do not (transitively) call __omp_parallel.  
    use(blockDim.x);  
}
```

```
__global__ void kernel() {  
    use(blockDim.x);  
}
```

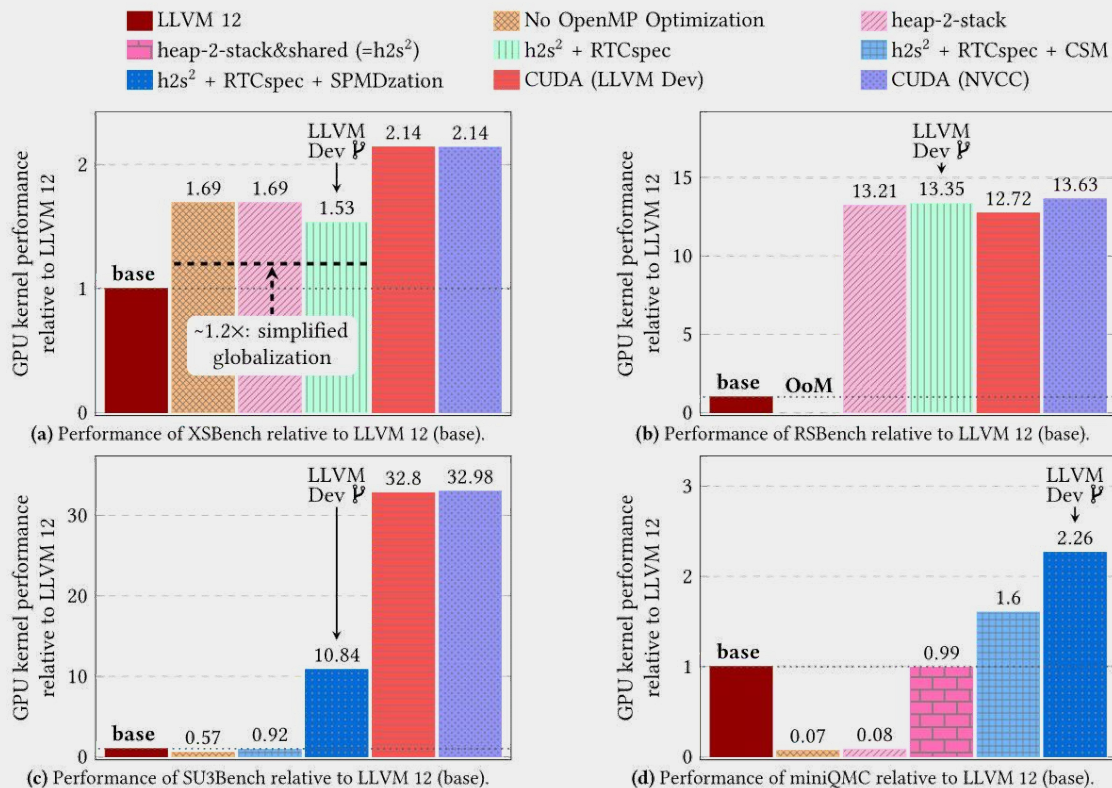


Simplifications, e.g., inlining, remove  
(now empty) abstraction layers.

⇒ CUDA-like code (IR and PTX)

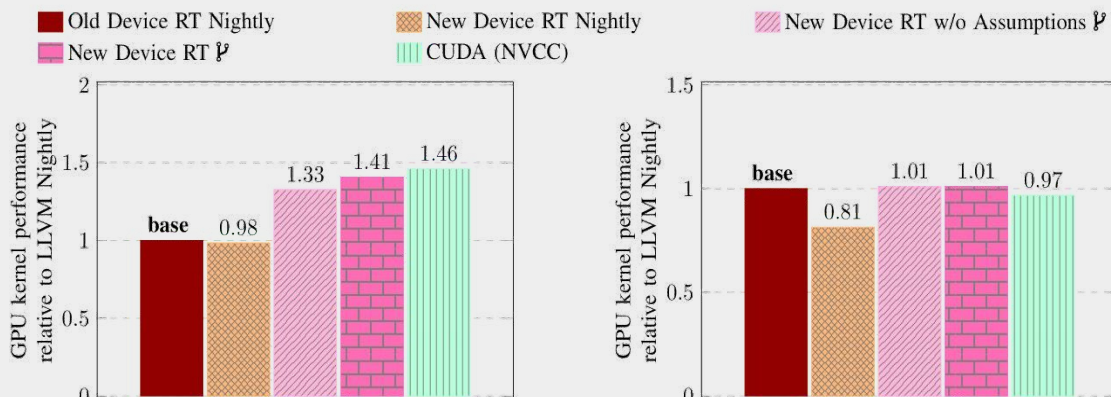
# Evaluation

# OpenMP Offloading Performance (Part I)







# OpenMP Offloading Performance (Part II)



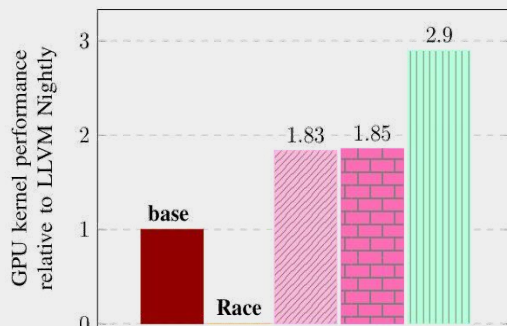
(a) Performance of XSbench relative to LLVM Nightly.



(b) Performance of RSBench relative to LLVM Nightly.



(c) Performance of TestSNAP relative to LLVM Nightly.



(d) Performance of MiniFMM relative to LLVM Nightly.

# Conclusion & Future Work

# Further Resources

- Official LLVM OpenMP documentation page
  - <https://openmp.llvm.org/>
- (OpenMP) Parallelism-Aware Optimizations (LLVM Dev'20)
  - <https://youtu.be/gtxWkeLCxmU>
- OpenMP Webinar ('20)
  - <https://www.openmp.org/events/webinar-a-compilers-view-of-the-openmp-api/>
- [Advancing OpenMP Offload Debugging Capabilities in LLVM](#) (LLPP'21)
- [Experience Report: Writing A Portable GPU Runtime with OpenMP 5.1](#) (IWOMP'21)

# Current and Future Work

- JIT and LTO for OpenMP offloading (see Lightning Talk)
- Profiling and runtime feature selection
- Heterogeneous host-device optimizations
- Generic GPU optimizations
- ...