

Otter: Tracing & Visualising OpenMP Programs as DAGs with LLVM OMPT

Adam Tuft

adam.s.tuft@durham.ac.uk

Task = self-contained unit of work + data

Create → Schedule → Do

- Creation: `task[loop]`
- Synchronisation: `taskgroup`, `taskwait`, `depend(...)`

```
int fibonacci(int n) {  
    int i, j;  
    if (n<2) return n;  
    #pragma omp task shared(i) firstprivate  
        (n)  
        i = f(n-1);  
    #pragma omp task shared(j) firstprivate  
        (n)  
        j = f(n-2);  
    #pragma omp taskwait  
    return i+j;  
}
```

The Problem

Want a "big picture" view of task-based code → **how?**





Figure 1: Performance analysis of a task-based Sudoku solver visualised in Vampir. Reproduced from [4].



Otter


A tool for visualising task creation and synchronisation

OMPT [1] + OTF2 [2] \rightarrow **Otter** [6]

- Visualise OpenMP task creation and synchronisation constructs
- Non-invasive
- Observe nested tasks & nested parallelism
- Callback-based (LLVM OMPT)
- OTF2 instrumentation layer

Example: fibonacci(5)

```
int fibonacci(int n) {  
    int i, j;  
    if (n<2) return n;  
    #pragma omp task shared(i)  
        firstprivate(n)  
        i = f(n-1);  
    #pragma omp task shared(j)  
        firstprivate(n)  
        j = f(n-2);  
    #pragma omp taskwait  
    return i+j;  
}
```



Example: 2D Euler Equations Solver




Figure 2: A single timestep in the 2D Euler solver, spawning background tasks for greater concurrency. Reproduced from [5].




Figure 3: A single timestep observed by Otter. Each task contains a domain traversal and background "enclave" tasks.

Tasking Inefficiencies Observed in LLVM RTL




Figure 4: Analysis of five time-steps with `OMP_NUM_THREADS=4`. Reproduced from [5].

- Number of queued tasks is capped at c. 1000 (4 threads).
- Pending tasks are processed immediately (unavailable for backfill).
- Need to be able to inform runtime to hold tasks back.

- API for data-driven taskification
- Measure work per task
- Support depend clause
- Detect critical path
- Relate to source
- Target tasks
- Device tracing
- Other tasking runtimes

Conclusion

- Can visualise task creation and synchronisation
- Provide data-driven insight into task-based code
- Unobstructed by particular task scheduling arrangement
- Need outstanding OMPT callbacks in LLVM

Acknowledgements

- Prof. Tobias Weinzierl
- Dr. Holger Schulz

Acknowledgements



The ExCALIBUR programme is supported by the UKRI Strategic Priorities Fund. The programme is led by the Met Office and the Engineering and Physical Sciences Research Council (EPSRC) along with the Public Sector Research Establishment, the UK Atomic Energy Authority (UKAEA) and UK Research and Innovation (UKRI) research councils, including the Natural Environment Research Council (NERC), the Medical Research Council (MRC) and the Science and Technologies Facilities Council (STFC).

Thank You!


References I

- [1] A. E. Eichenberger et al. “OMPT: An OpenMP Tools Application Programming Interface for Performance Analysis”. In: *OpenMP in the Era of Low Power Devices and Accelerators*. Ed. by A. P. Rendell, B. M. Chapman, and M. S. Müller. Berlin, Heidelberg: Springer Berlin Heidelberg, 2013, pp. 171–185. ISBN: 978-3-642-40698-0. DOI: 10.1007/978-3-642-40698-0_13.
- [2] D. Eschweiler et al. “Open Trace Format 2: The Next Generation of Scalable Trace Formats and Support Libraries”. In: vol. 22. Jan. 2012, pp. 481–490. ISBN: 9781614990406. DOI: 10.3233/978-1-61499-041-3-481.
- [3] OpenMP Architecture Review Board. *OpenMP Application Program Interface Version 5.0*. Nov. 2018. URL: <https://www.openmp.org/wp-content/uploads/OpenMP-API-Specification-5.0.pdf>.

- [4] D. Schmidl et al. “Suitability of Performance Tools for OpenMP Task-Parallel Programs”. In: *Tools for High Performance Computing 2013*. Ed. by A. Knüpfer et al. Cham: Springer International Publishing, 2014, pp. 25–37. ISBN: 978-3-319-08144-1. DOI: 10.1007/978-3-319-08144-1_3.
- [5] H. Schulz et al. “Task Inefficiency Patterns For A Wave Equation Solver”. In: *arXiv preprint arXiv:2105.12739* (2021). URL: <https://arxiv.org/abs/2105.12739>.
- [6] A. Tuft. *Otter v0.1*. Version v0.1-submit. Aug. 2021. DOI: 10.5281/zenodo.5259058. URL: <https://doi.org/10.5281/zenodo.5259058>.


Example: Single Taskloop

```
#pragma omp parallel num_threads(2)
{
    #pragma omp single nowait
    #pragma omp taskloop nogroup
    for (int j=0; j<4; j++)
    {
        // do work
    }
}
```




Post-processing Algorithm

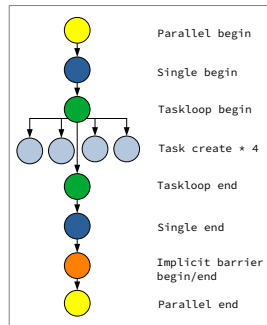
a) Events are grouped in (possibly nested) chunks corresponding to a thread's view of a parallel, single, master or task region.



b) Chunks generate sub-graphs of the final graph. Labels indicate the thread-event pair represented. Dashed black lines show chunk borders and red lines connect nodes representing the same OpenMP construct as viewed by different threads.



c) The disjoint union of the sub-graphs in (b) produces the final graph with equivalent nodes and chunk borders collapsed, and nodes coloured by the type of OpenMP construct they represent.



Allowed values of `ompt_set_callback`

Table 4.2 of [3]:

Return code abbreviation	N	S/P	A	Return code abbreviation	N	S/P	A
<code>ompt_callback_thread_begin</code>			Y	<code>ompt_callback_mutex_released</code>	Y	Y	Y
<code>ompt_callback_thread_end</code>			Y	<code>ompt_callback_dependences</code>	Y	Y	Y
<code>ompt_callback_parallel_begin</code>			Y	<code>ompt_callback_task_dependence</code>	Y	Y	Y
<code>ompt_callback_parallel_end</code>			Y	<code>ompt_callback_work</code>	Y	Y	Y
<code>ompt_callback_task_create</code>			Y	<code>ompt_callback_master</code>	Y	Y	Y
<code>ompt_callback_task_schedule</code>			Y	<code>ompt_callback_target_map</code>	Y	Y	Y
<code>ompt_callback_implicit_task</code>			Y	<code>ompt_callback_sync_region</code>	Y	Y	Y
<code>ompt_callback_target</code>			Y	<code>ompt_callback_reduction</code>	Y	Y	Y
<code>ompt_callback_target_data_op</code>			Y	<code>ompt_callback_lock_init</code>	Y	Y	Y
<code>ompt_callback_target_submit</code>			Y	<code>ompt_callback_lock_destroy</code>	Y	Y	Y
<code>ompt_callback_control_tool</code>			Y	<code>ompt_callback_mutex_acquire</code>	Y	Y	Y
<code>ompt_callback_device_initialize</code>			Y	<code>ompt_callback_mutex_acquired</code>	Y	Y	Y
<code>ompt_callback_device_finalize</code>			Y	<code>ompt_callback_nest_lock</code>	Y	Y	Y
<code>ompt_callback_device_load</code>			Y	<code>ompt_callback_flush</code>	Y	Y	Y
<code>ompt_callback_device_unload</code>			Y	<code>ompt_callback_cancel</code>	Y	Y	Y
<code>ompt_callback_sync_region_wait</code>	Y	Y	Y	<code>ompt_callback_dispatch</code>	Y	Y	Y

N = `ompt_set_never`, S = `ompt_set_sometimes`, P = `ompt_set_sometimes_paired`, A = `ompt_set_always`

Region Symbols



parallel



implicit barrier



taskloop



initial task



taskwait



single



explicit task



taskgroup



for-loop



master

Figure 5: Node style encodes the type of OpenMP construct.

main -
3 branches
1 tag

Go to file
Add file -
Code -

adamtuft Update README.md

cd8f417 · 22 hours ago · 328 commits

docs	Add symbol table to readme	yesterday
include	Minor corrections, remove some redundant code, fix python-igraph vers...	22 hours ago
lib	Set up project template	4 months ago
obj	Set up project template	4 months ago
src	Minor corrections, remove some redundant code, fix python-igraph vers...	22 hours ago
.gitignore	Update readme	yesterday
LICENSE	Create LICENSE	2 months ago
Makefile	Remove redundant targets	22 hours ago
README.md	Update README.md	22 hours ago
otter-defs.sh	Remove redundant otter #define's, slight update to Makefile	2 months ago
otter-env.sh	Correct names of environment variables	22 hours ago

README.md

Otter - An OMPT Tool for Tracing OpenMP Tasks

DOI: 10.5281/zenodo.5259058

Otter is a tool for visualising the structure of task-based OpenMP programs allowing developers and researchers to see the true structure of their OpenMP 5.0 programs from the perspective of the OpenMP runtime, without any modification of the target application.

Otter uses the OpenMP Tools interface in [OpenMP 5.0](#) to observe task creation and synchronisation events, extracting from this data the structure of a target application independent of the particular scheduling of tasks at runtime.

About

OMPT Task Tracer

Readme

BSD-3-Clause License

Releases 1

Initial Release - MSc Submi... Latest

22 hours ago

Packages

No packages published

[Publish your first package](#)

Languages

C 77.1%

Python 19.8%

Makefile 1.6%

C++ 1.1%

Shell 0.4%

github.com/adamtuft/otter